

# # Basic File Systems

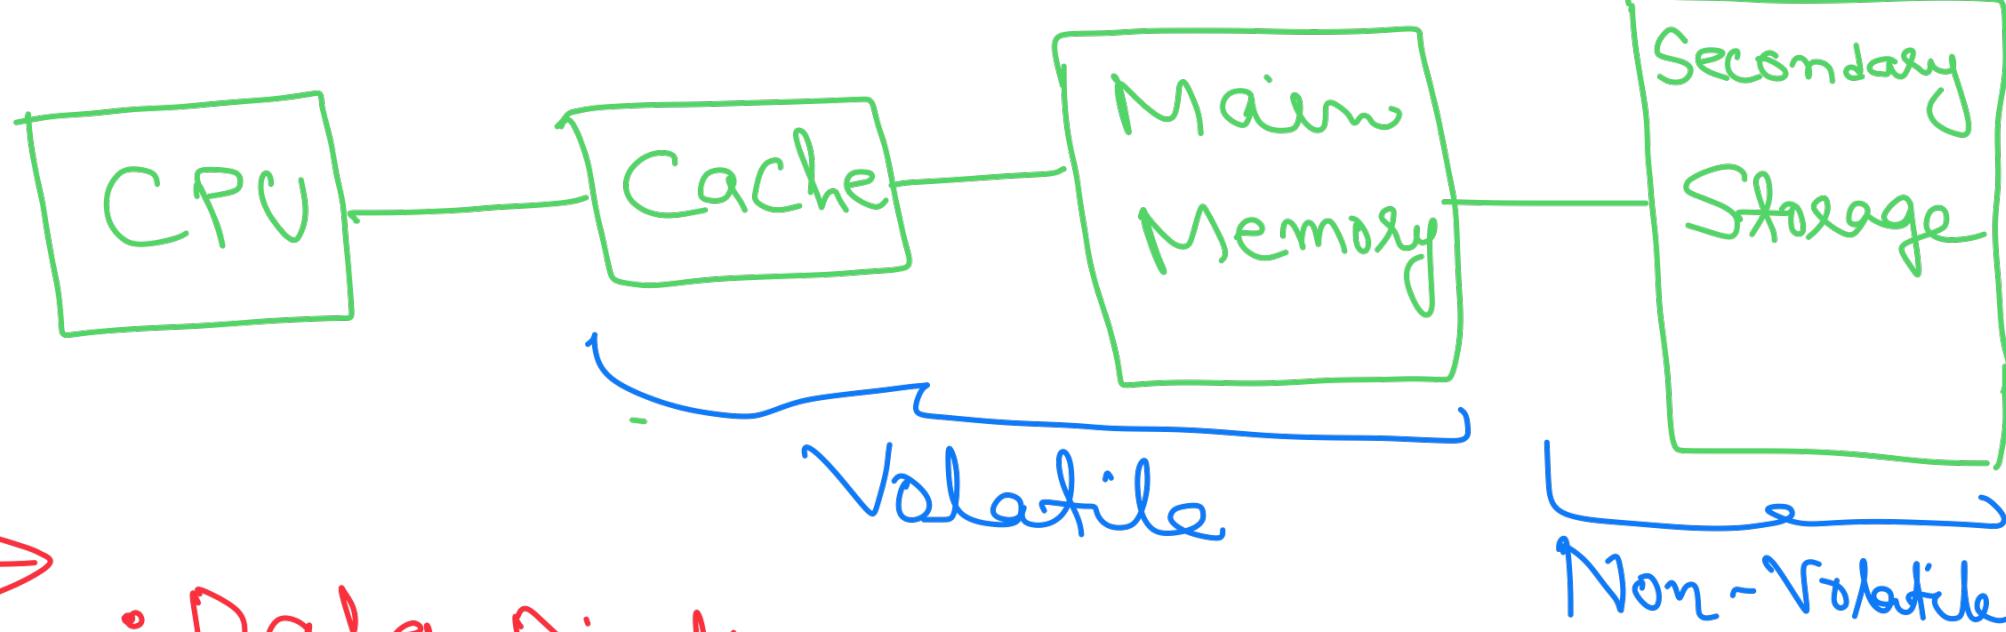
---

- Data Dictionary
- Database
- Intermediate Computation
- Queries

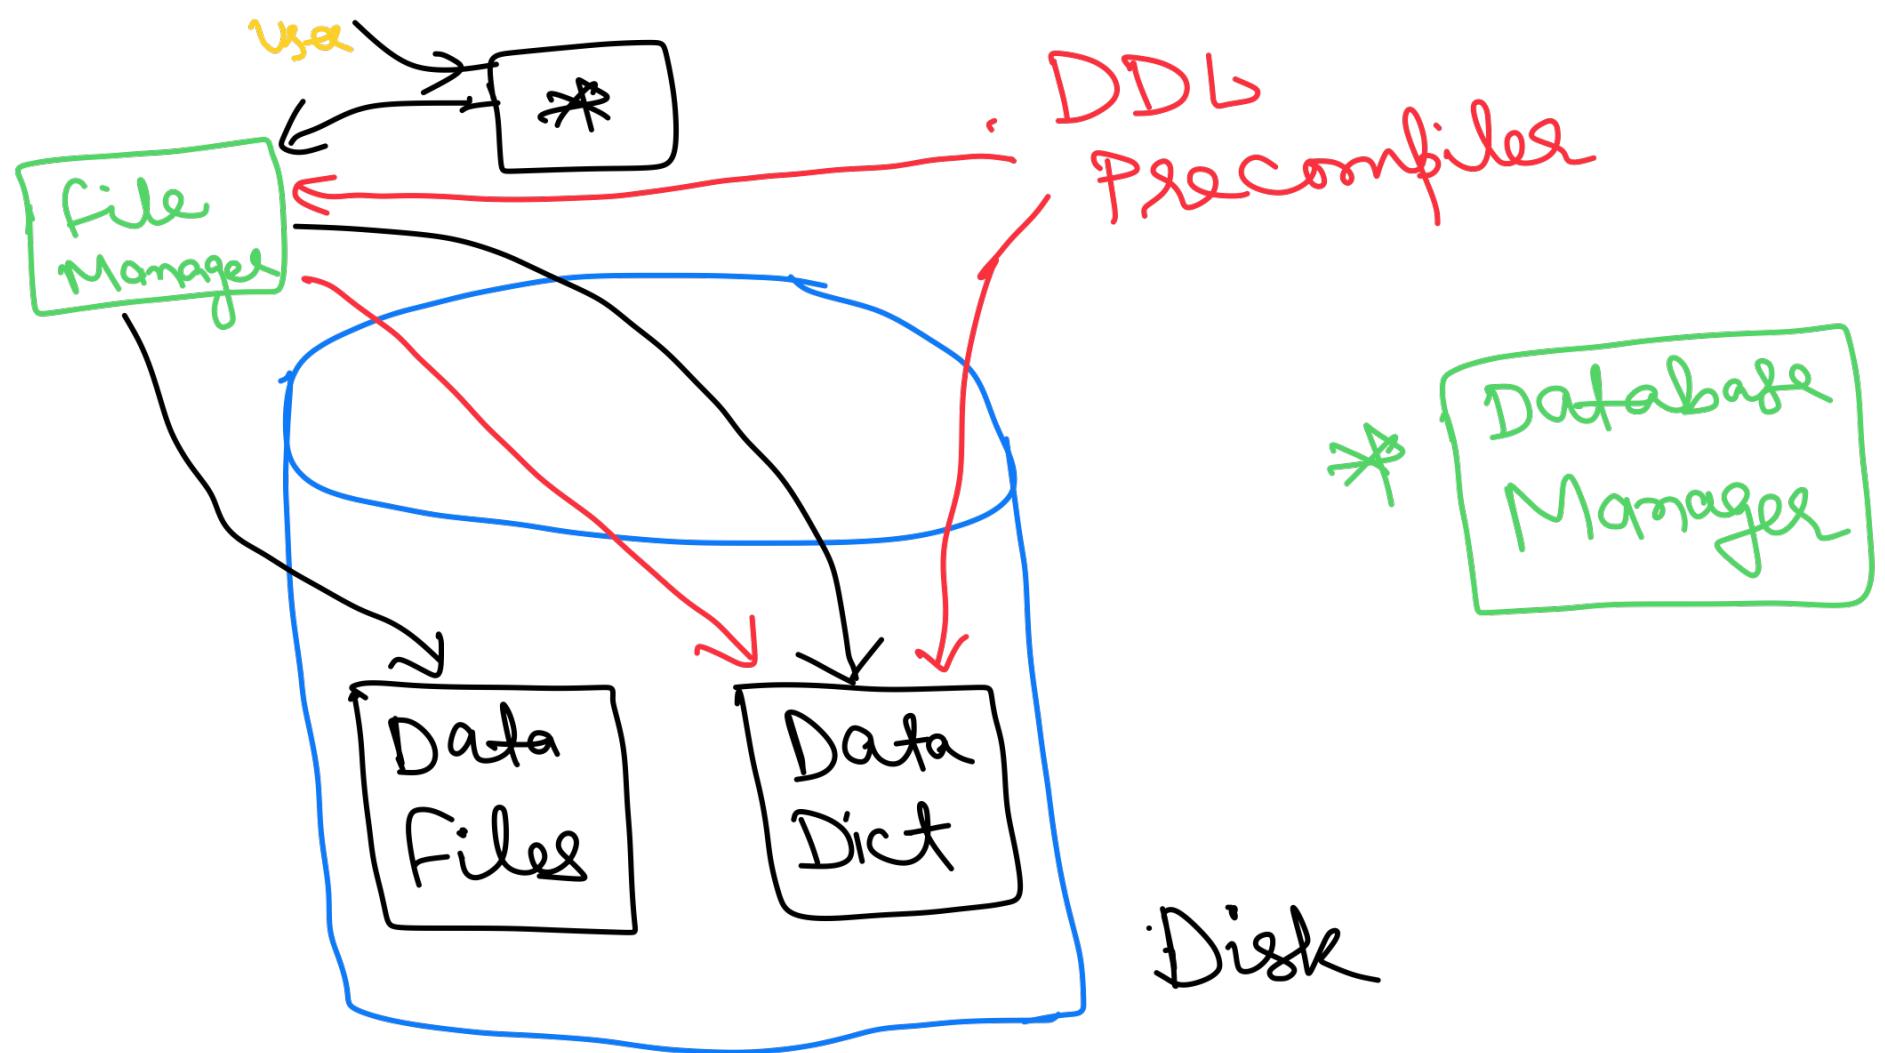
# # Physical Storage Media

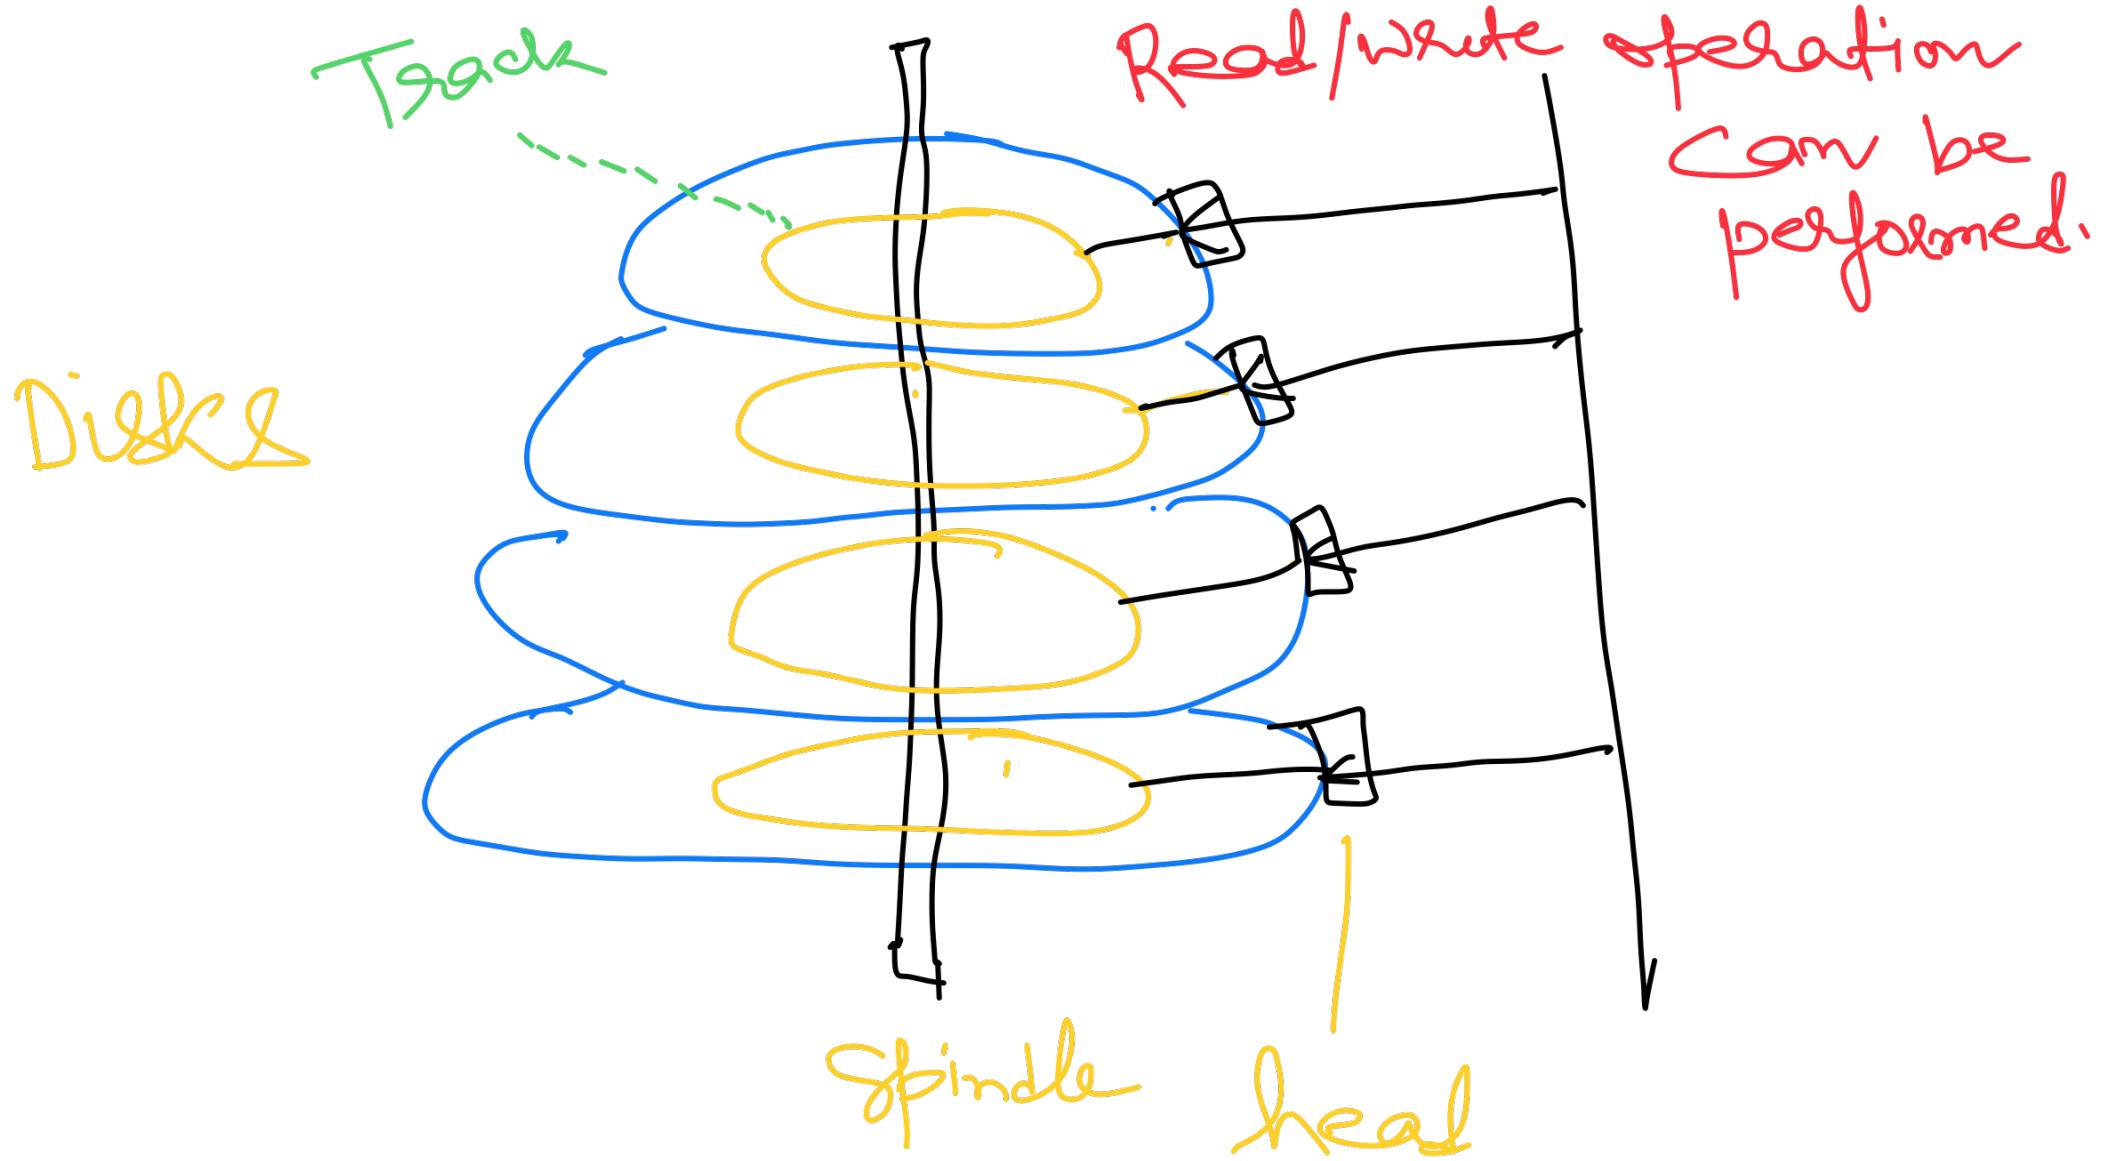
---

- Cache Memory
- Main Memory
- Disk Storage
- Tape Storage



- Data Dictionary
  - Database
  - Intermediate Computation
  - Queries
- Non-volatile
- Can/May  
be  
performed  
on Main  
Memory





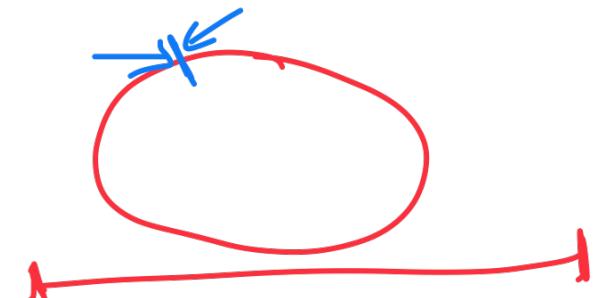
## \* Disk Storage

- Seek Time
  - Latency
- Latency → Time required to place the head on the Track
- Latency → Time required to Read some unit of data.

## # Disk Organization

- Track
- Blocks
- Records
- Files

(Maybe Cylinder which contain multiple tracks)



→ In hard disk, Read / Write operation happened blockwise.

However, in main memory, Read / Write operation happened byte wise.

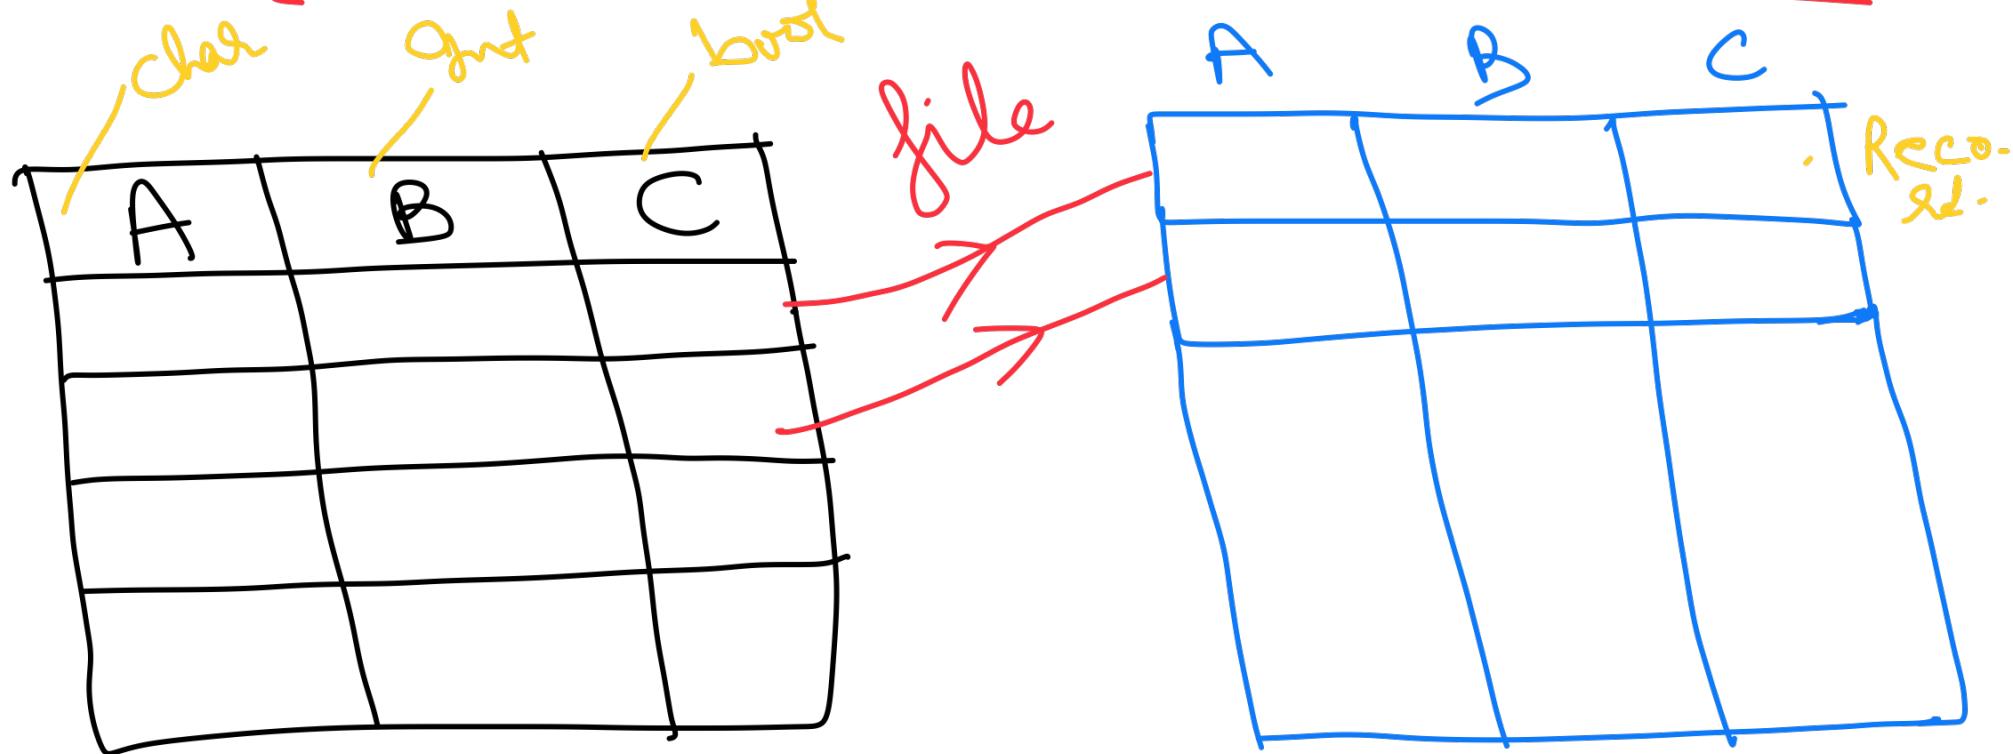
One block contains multiple bytes.

## # Records in Blocks

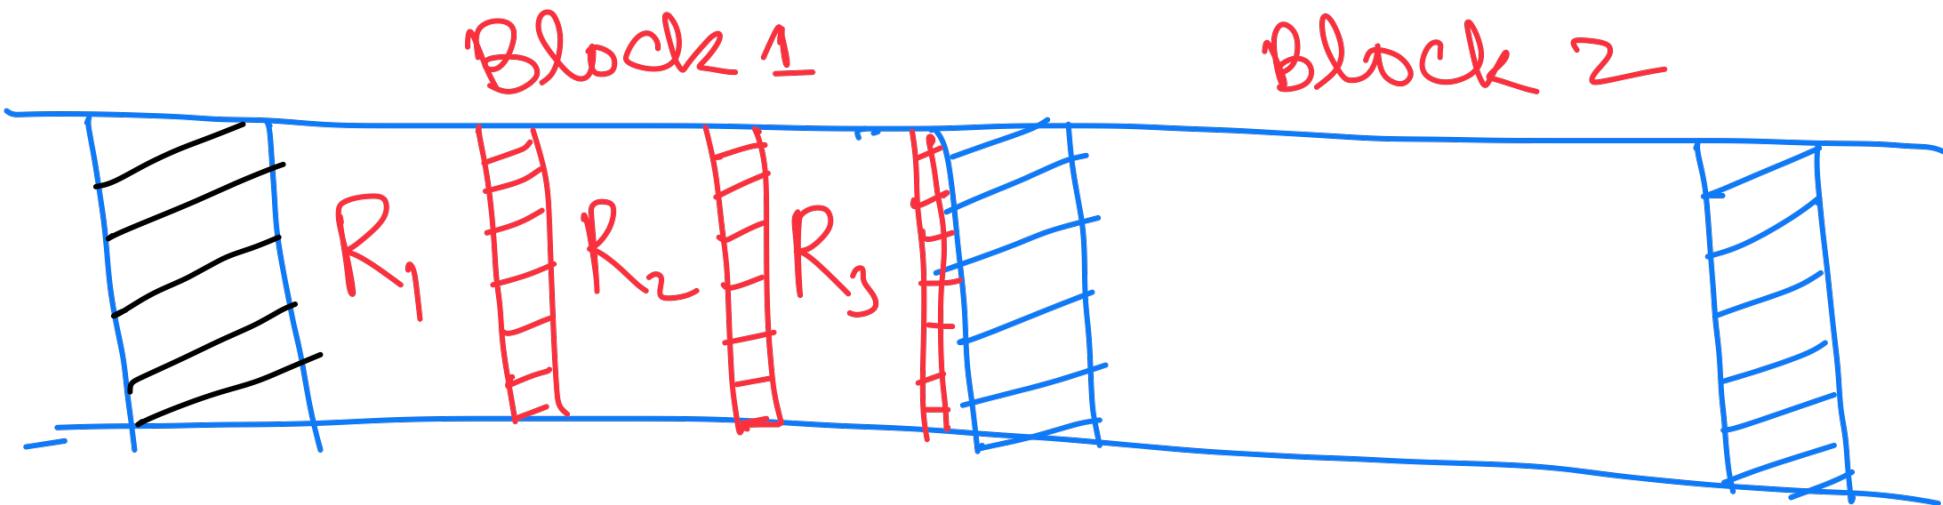
---

- \* fixed length blocking for fixed length records
- \* Variable length Unspanned Records
- \* Variable length Spanned Records

# # Fixed Length Records

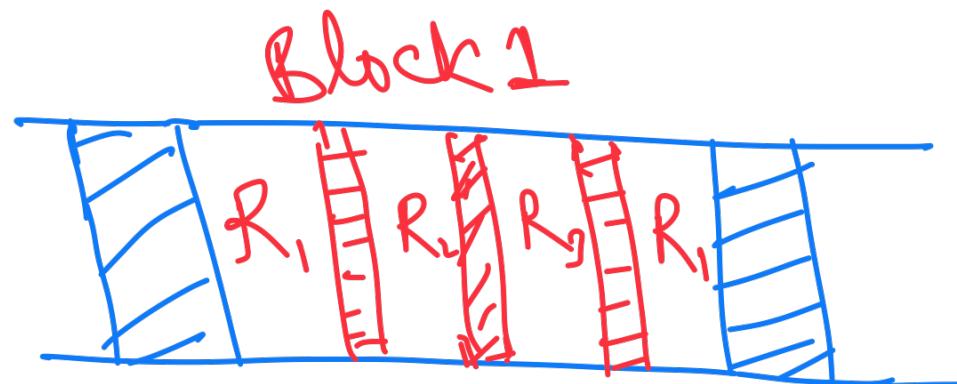
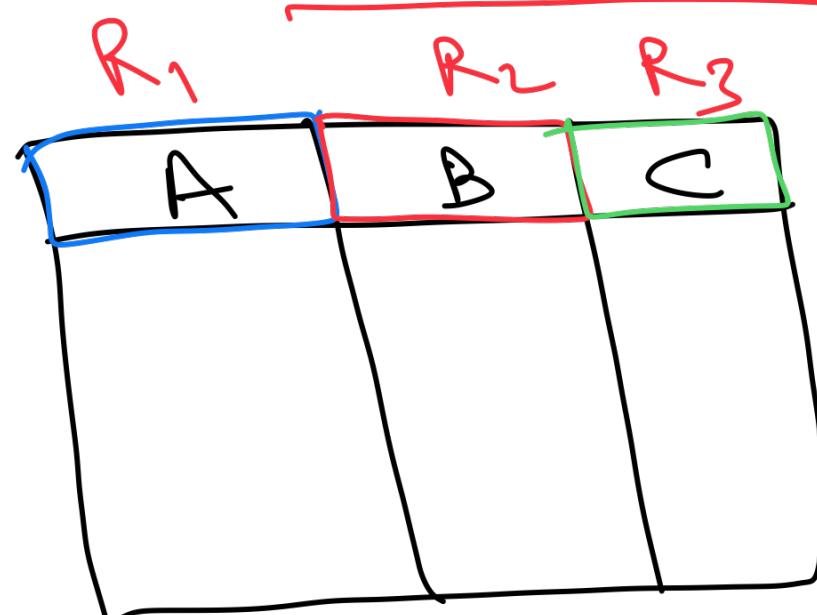


One Table per file



Block 2

# Variable Length Records (VLR)



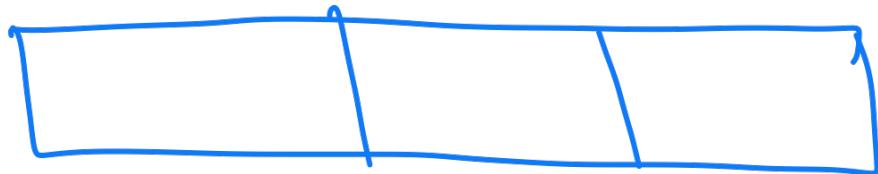
A	B	C

In one file,  
we may store  
more than one  
table

- Unspanned type VLR - gt  
does not allow same record to be part of more than one block.
- Spanned type VLR - gt allows.

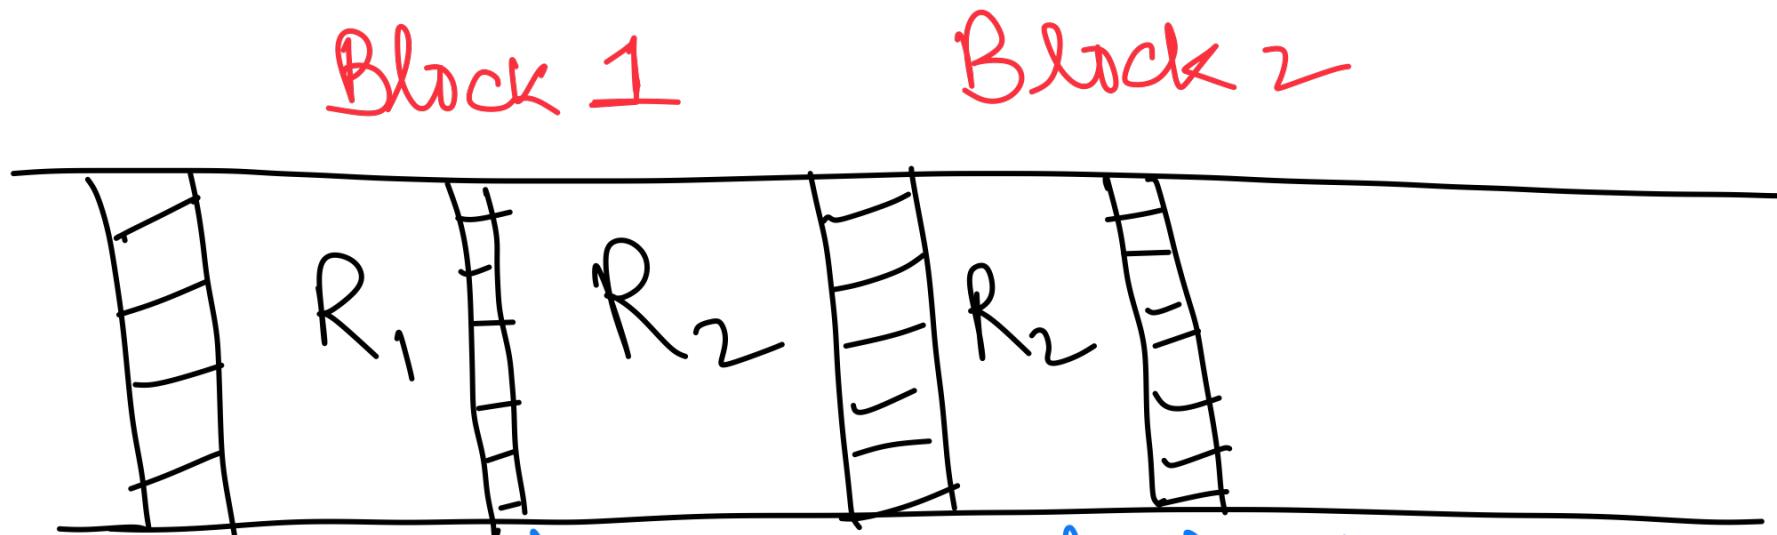
ACC_NO	YR_PUB	TITLE

TITLE → YR\_PUB ACC\_NO



⋮

# # Variable length Spanned Records



We allow part of record to be in one block and rest of it to be on the next block.

# # Pinned and Unpinned Records

## # Files

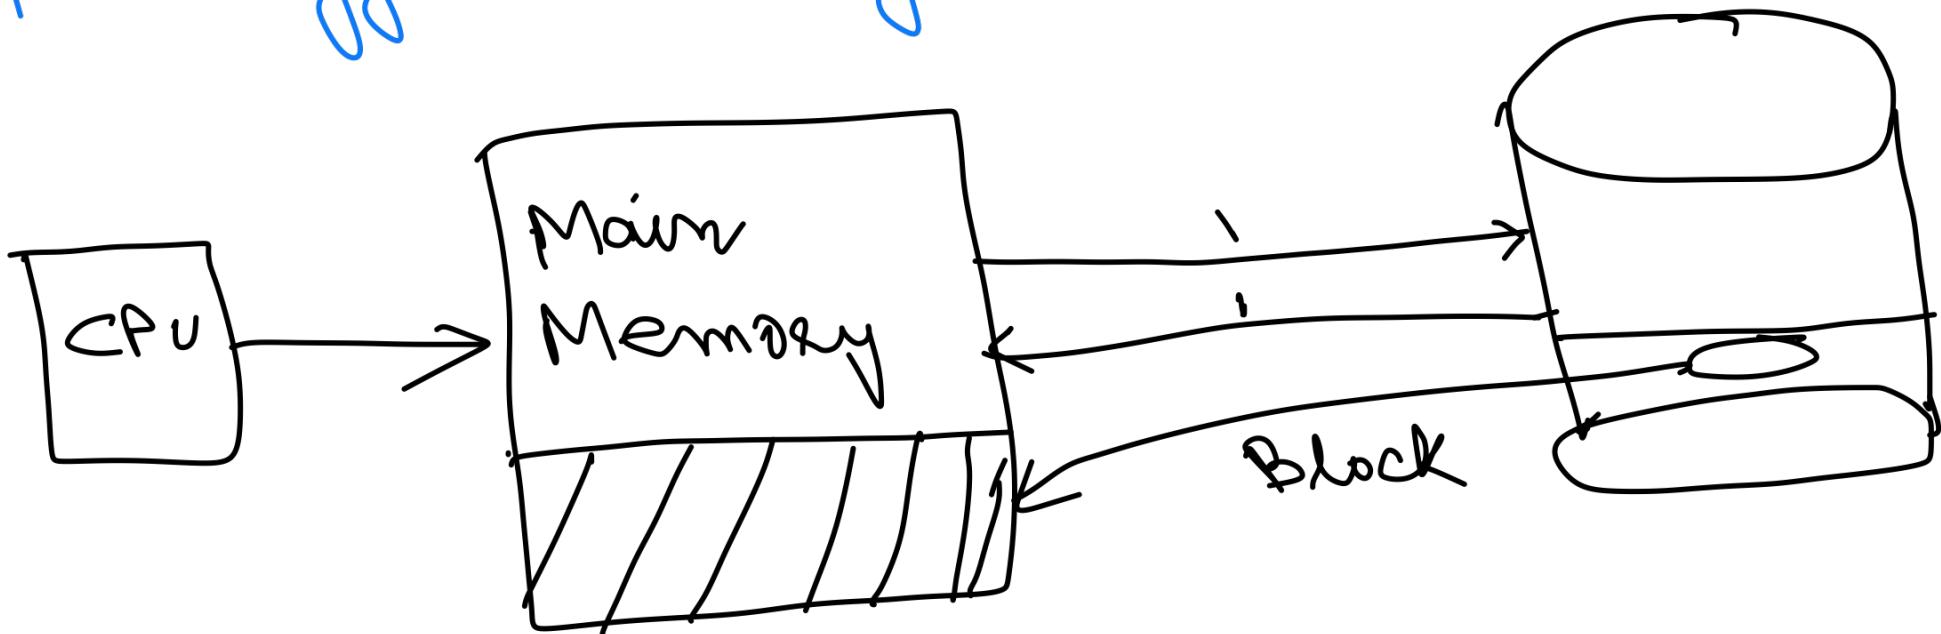
### \* Data Dictionaries

- Relation Names
- Attribute Names, Domains
- View Definitions
- Integrity Constraints
- Access Rights
- Corresponding Data files

## \* Data files

- One relation one file
- One relation many files
- Many relation one file

## \* Buffer Management



→ Locality of Reference

→ Block Replacement Policy

→ Pinned and Unpinned Records

## # File Organization

- Heap or file
  - Sequential files
  - Indexed Sequential files
  - Hashed files
  - B+ - Trees
- A file refers to a collection of related records, i.e. data, which are stored and organized logically.

# # Heap File (Pile)



Very costly

Insertion : By Stacking

Deletion : By setting a field in Record

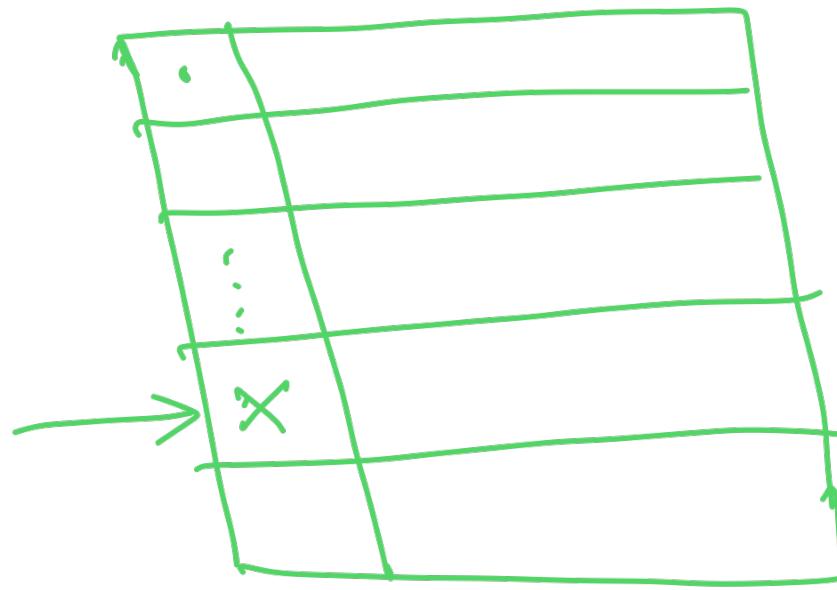
Search : Brute Force

Reorganize the file  
(Remove the Deleted entries)

Garbage Collection

## # Sequential files

Records are kept in sequential order  
on some attributes (s) - usually the  
Primary Key



5721	Compiler Design	1978
5732	Database Engg	1986
5734	Database Engg	1986
5735	Algorithms	1992
5740	Graphics	1997

5733	Graph Theory	1976
------	--------------	------

Insert  
 Delete  
 Block Search → Search | Find

Block Search

Gets a block

Binary Search in the block

\* Acc-No Title YR-PUB NEXT

Acc-No	Title	YR-PUB	NEXT
5721	Compiler Del.	1978	
5732	Database	1986	
5734	Database	1986	
5735	Algorithms	1992	
5740	Graphics	1994	771
5733	Graphics	1976	

Insert

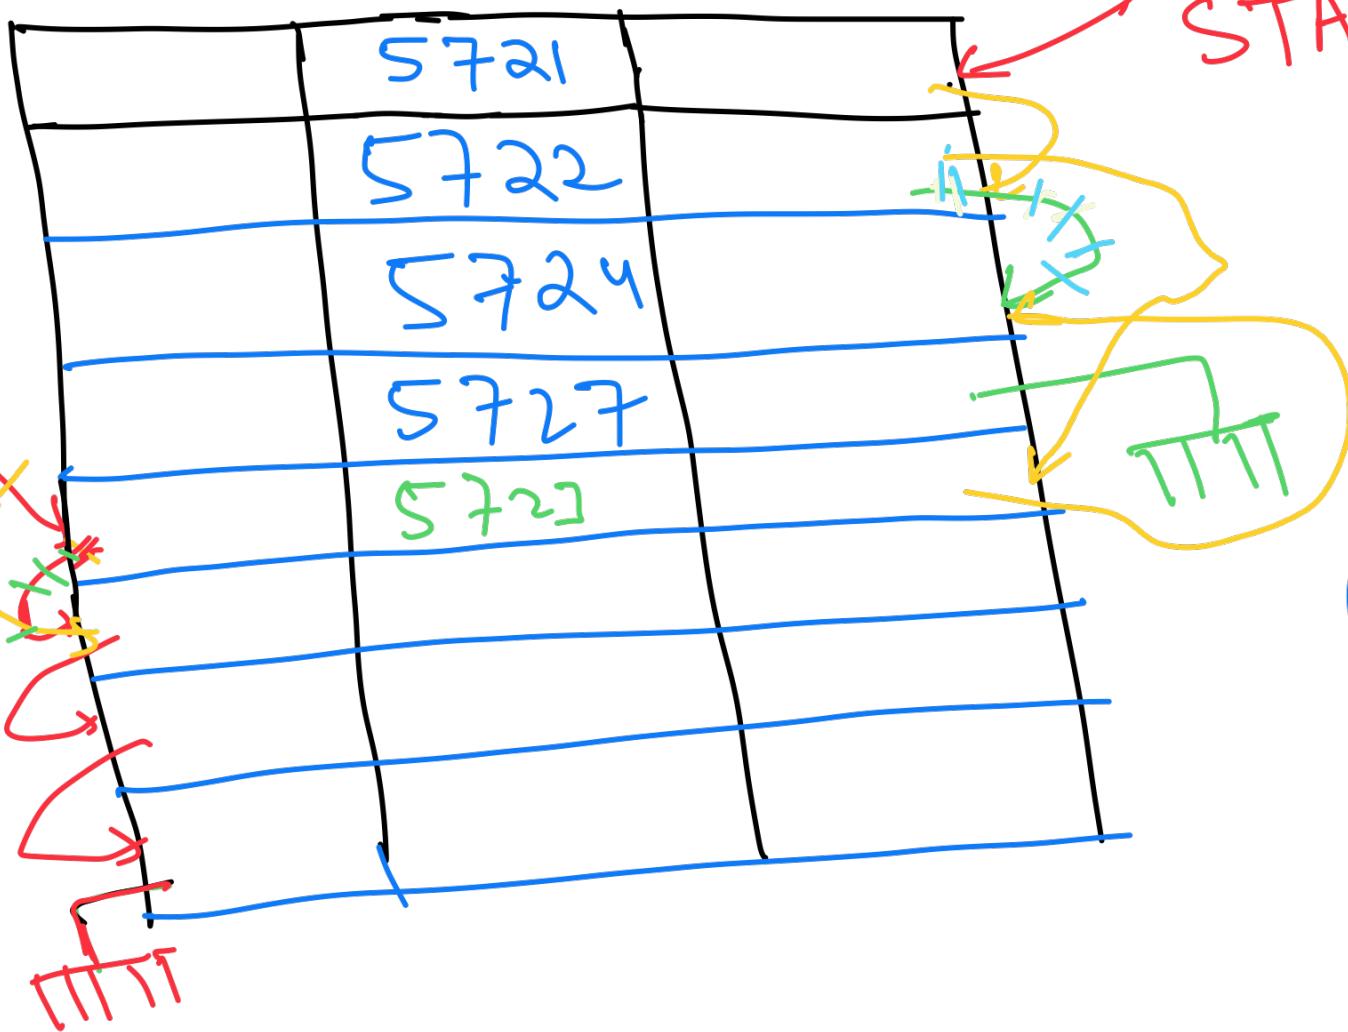
The diagram illustrates the insertion of a new record into a linked list. A green arrow points to the row for book 5734, indicating the insertion point. A yellow box highlights the new record: Acc-No 5733, Title Graphics, YR-PUB 1976, and NEXT 771. Yellow arrows show the connections between the old records and the new one, and another arrow points from the new record back to the original list.

AVB

ACC-No	Title	YR_PUB	Next
5721	Compiler Des.	1978	
5732	Database	1986	
5734	Database	1986	
5735	Algorithms	1992	
5740	Graphics	1994	777

Reorganization

AVL

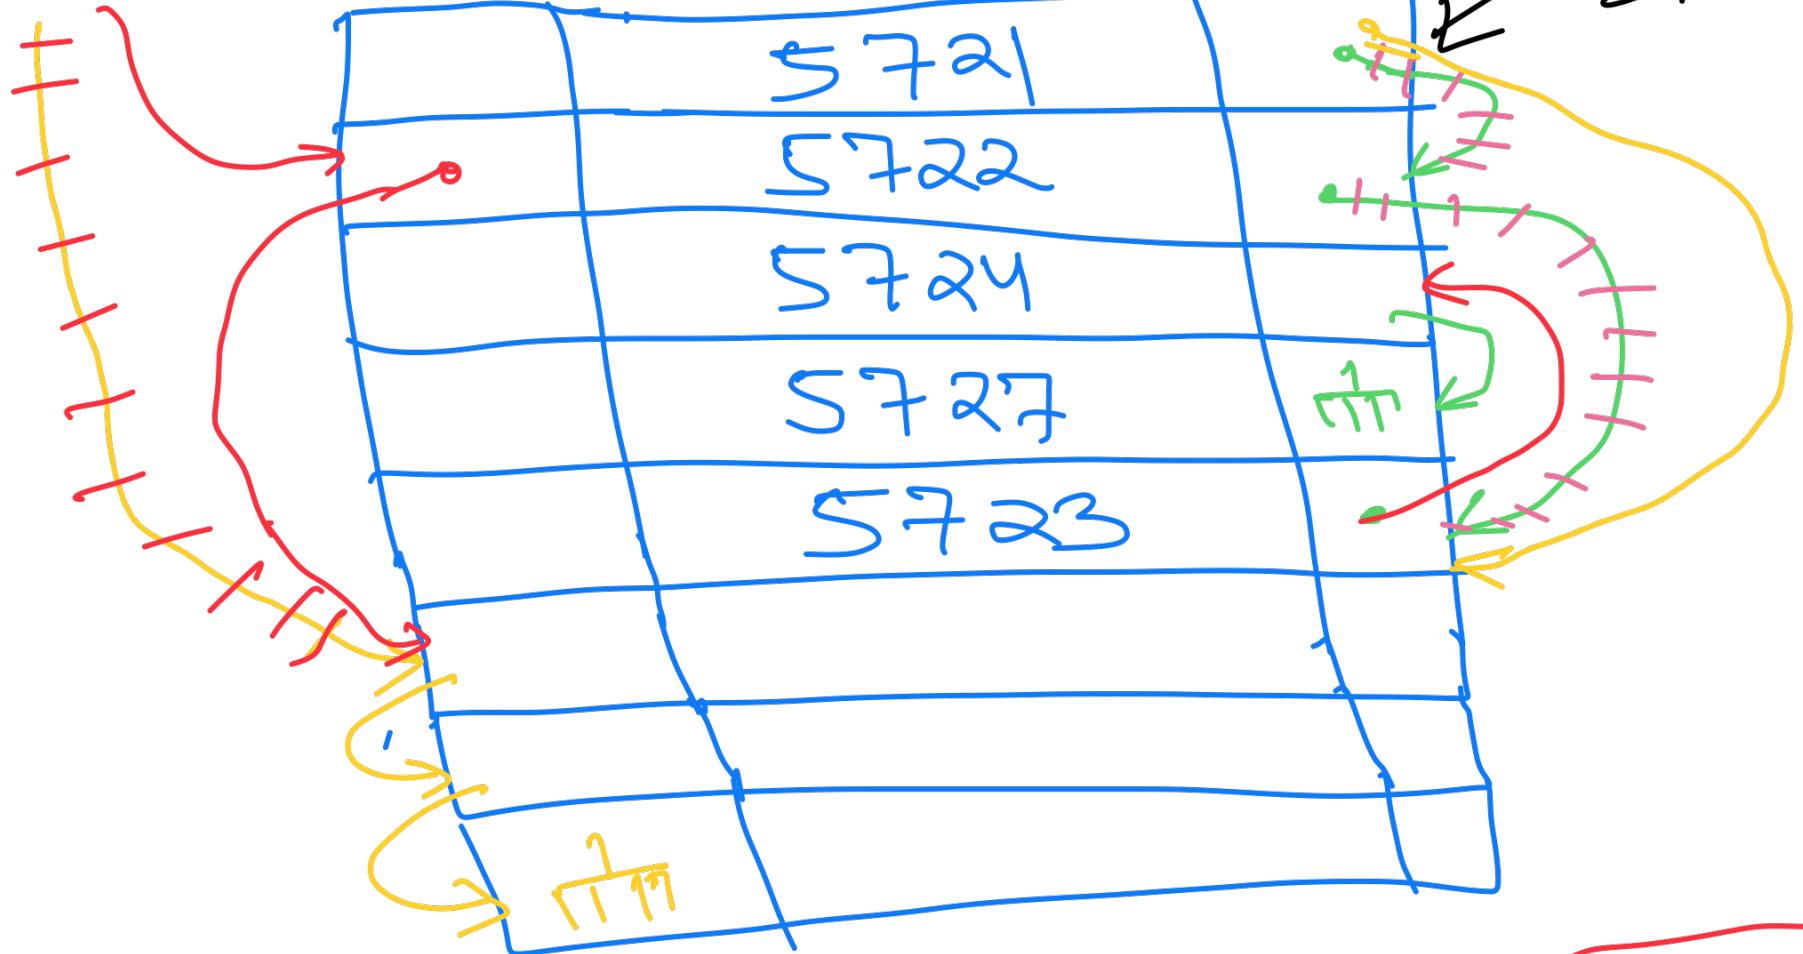


START

Insert

5723

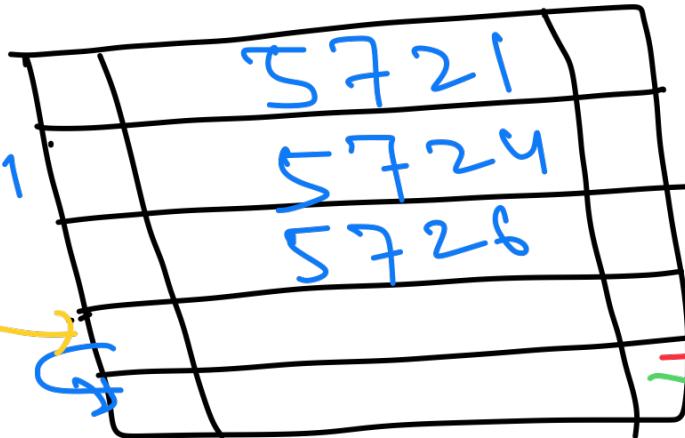
Avb



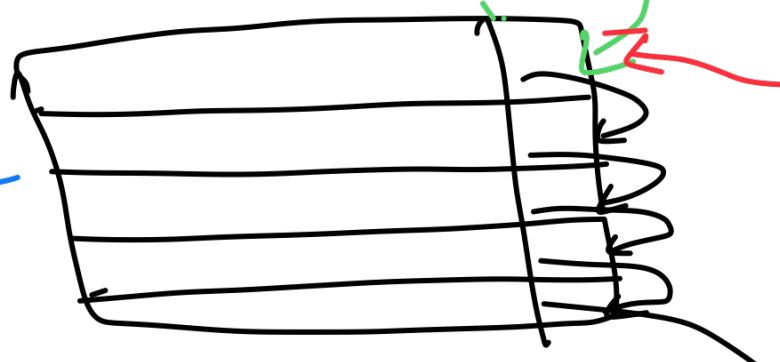
Delete  
5722

AVL

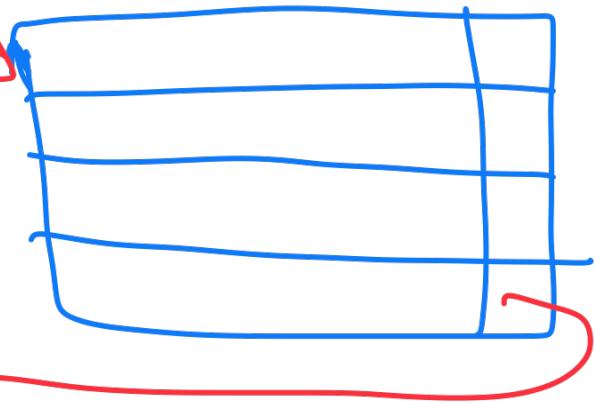
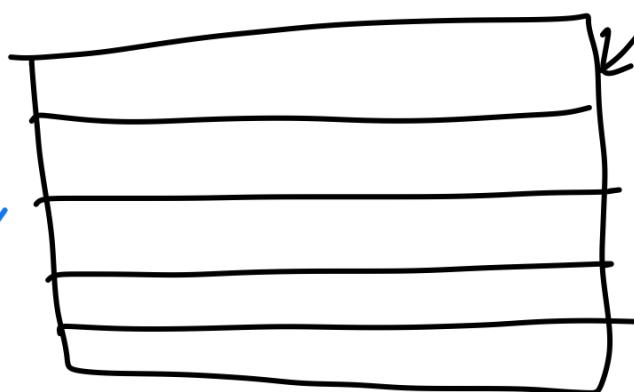
Block 1



Block ~



Block n



Reorganize  
the "block"

→ Sequential access in sequential file, hence it is making process slow.

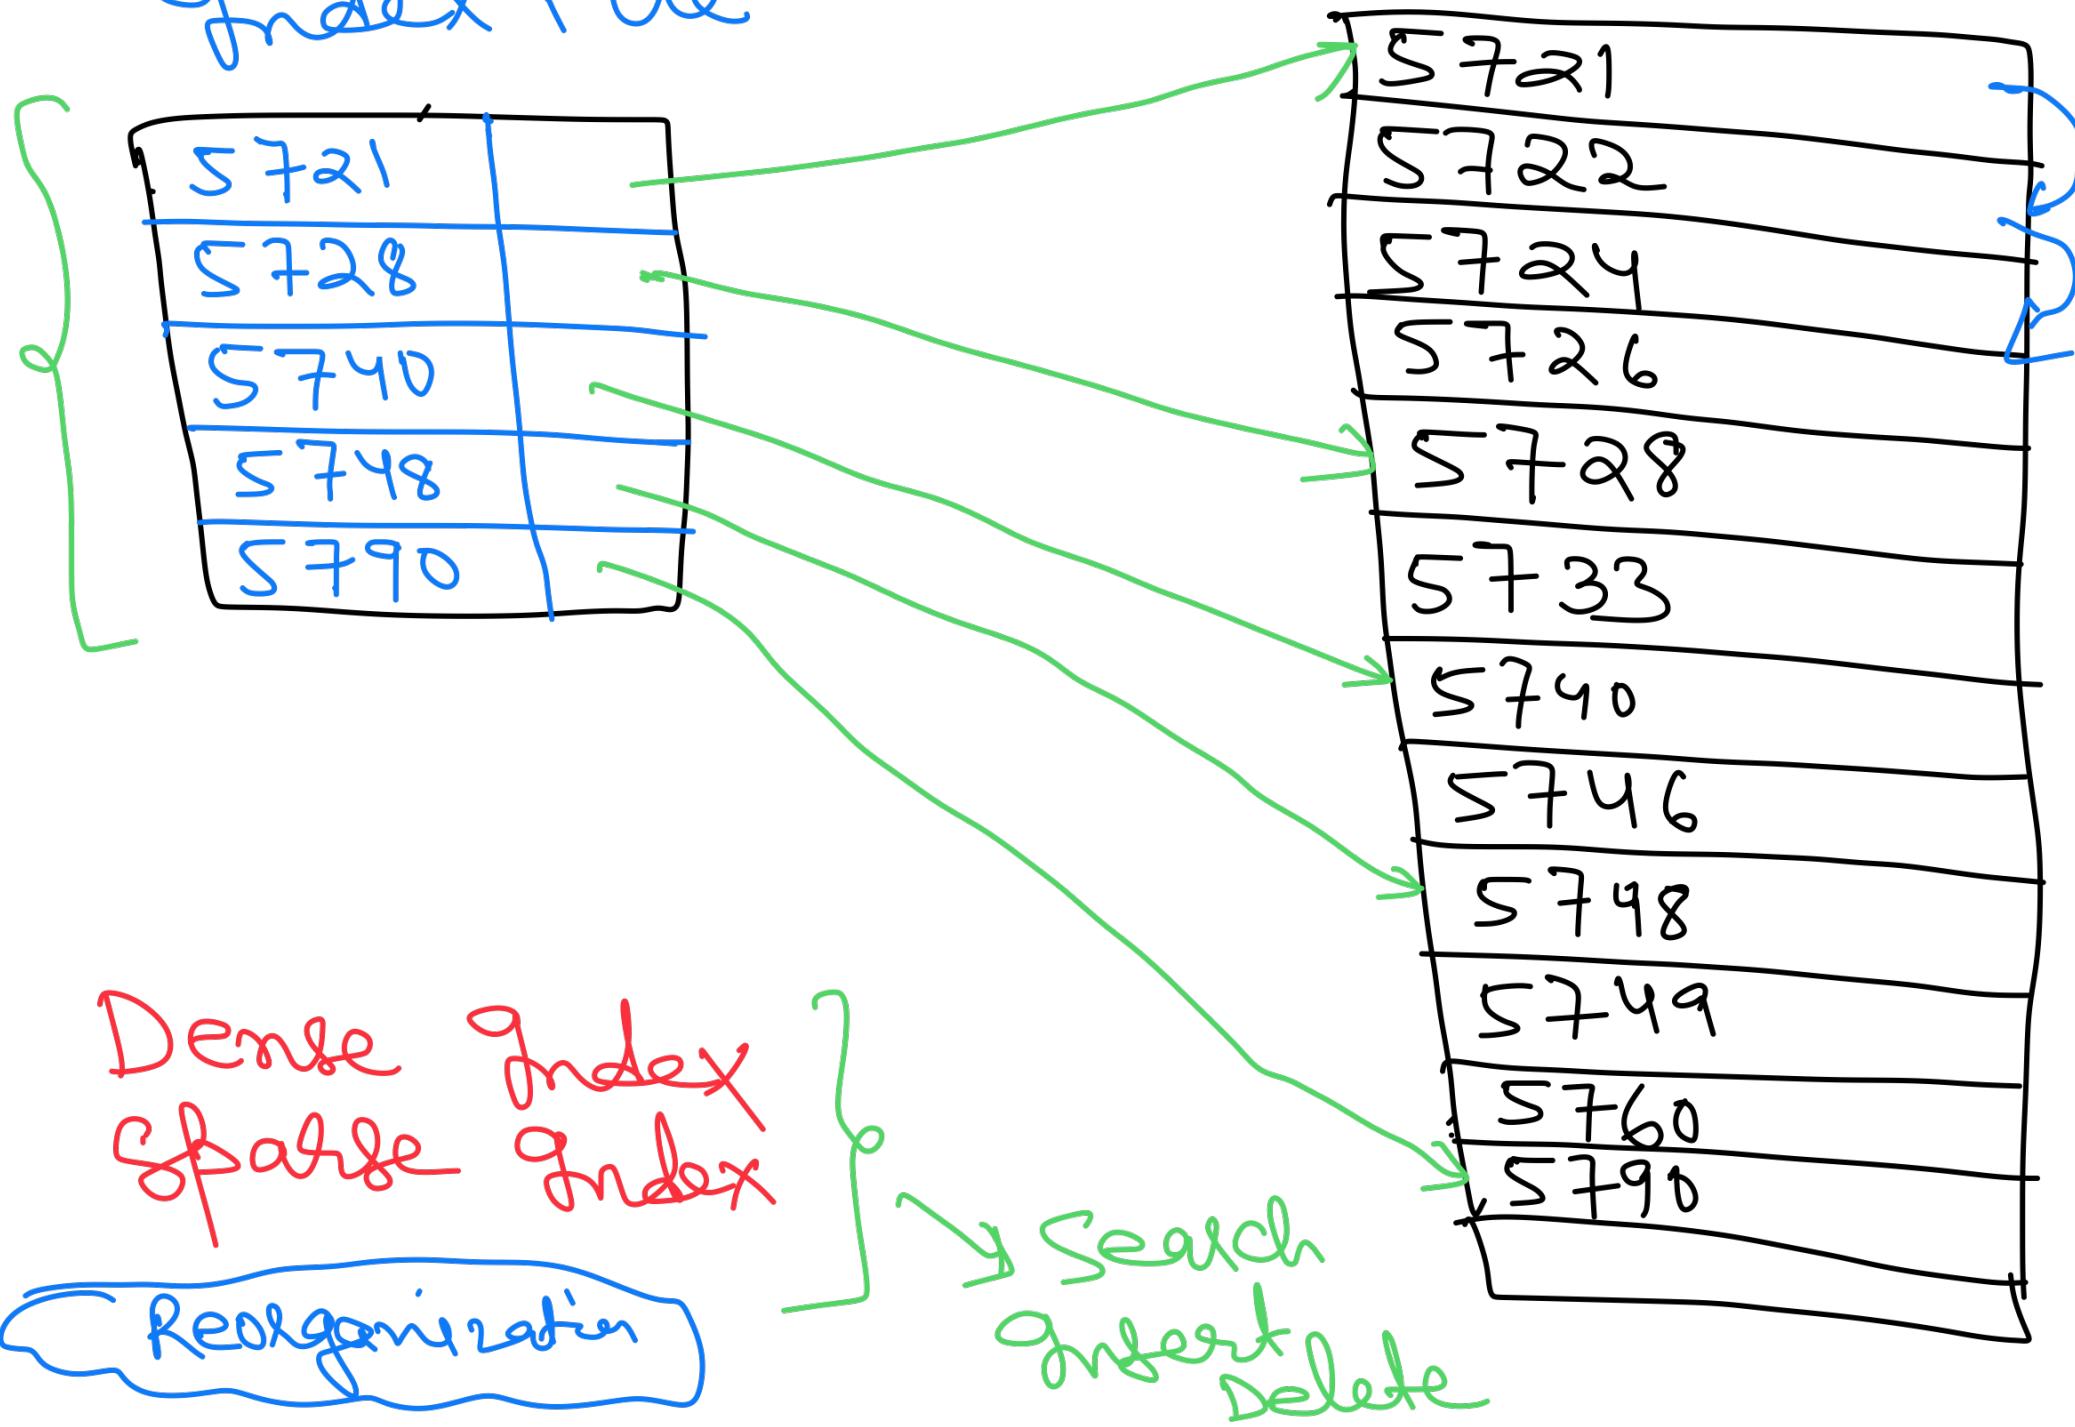
## # Indexed Sequential Files

Maintain an index over the SEQ file -

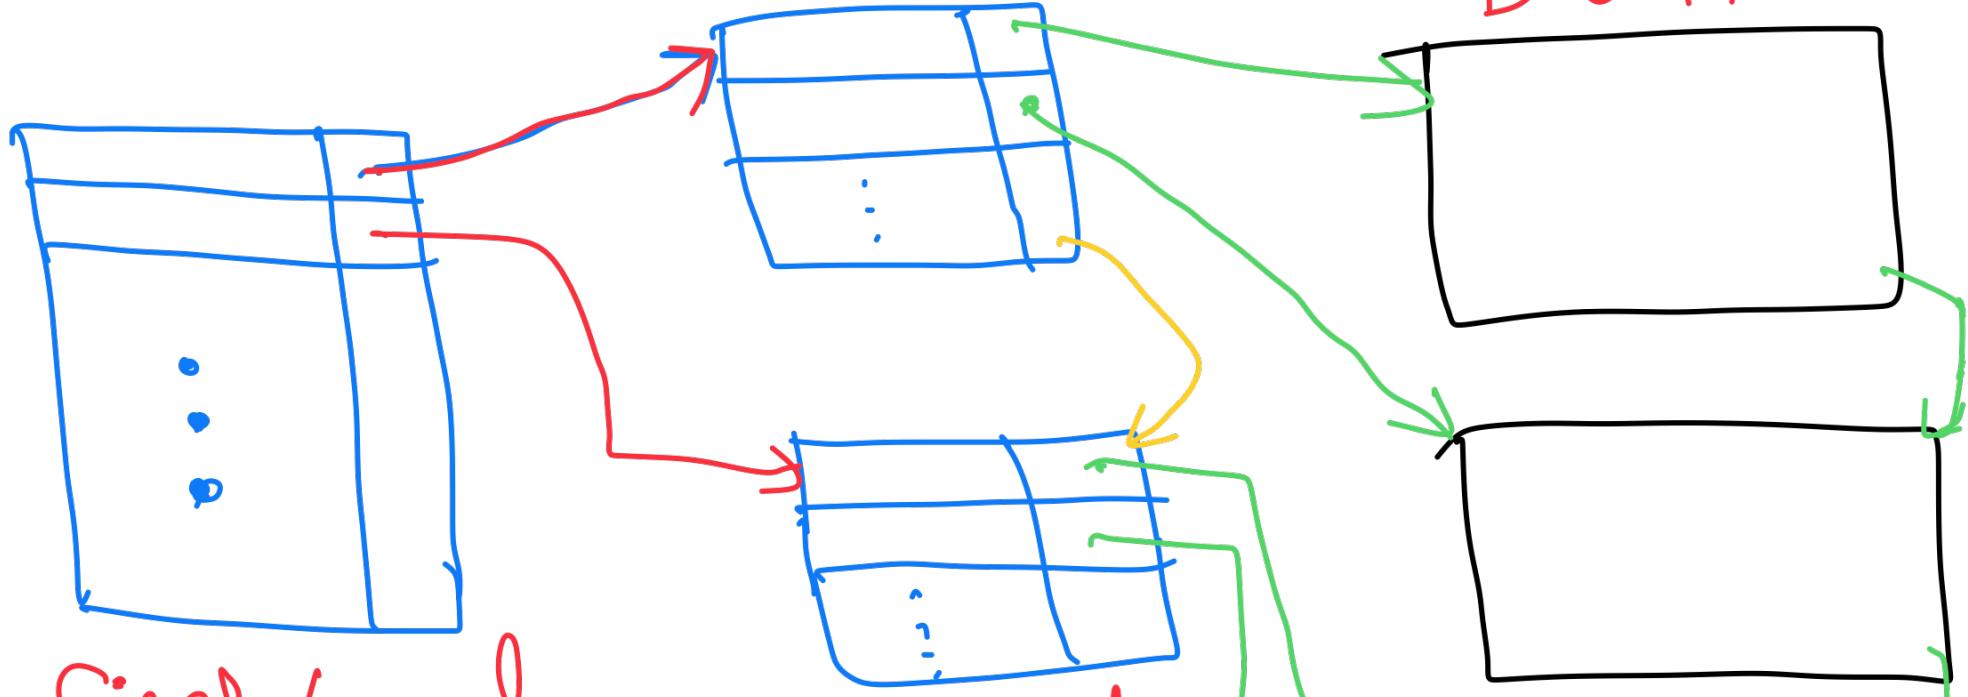


Smaller file

# gindex file

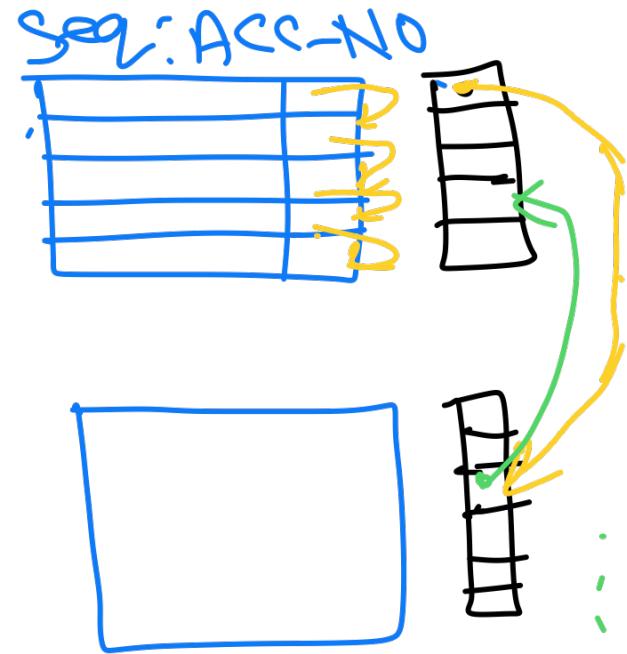


Data File



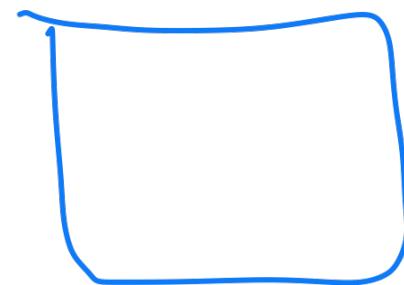
Multi-Level Sparse  
Indexing

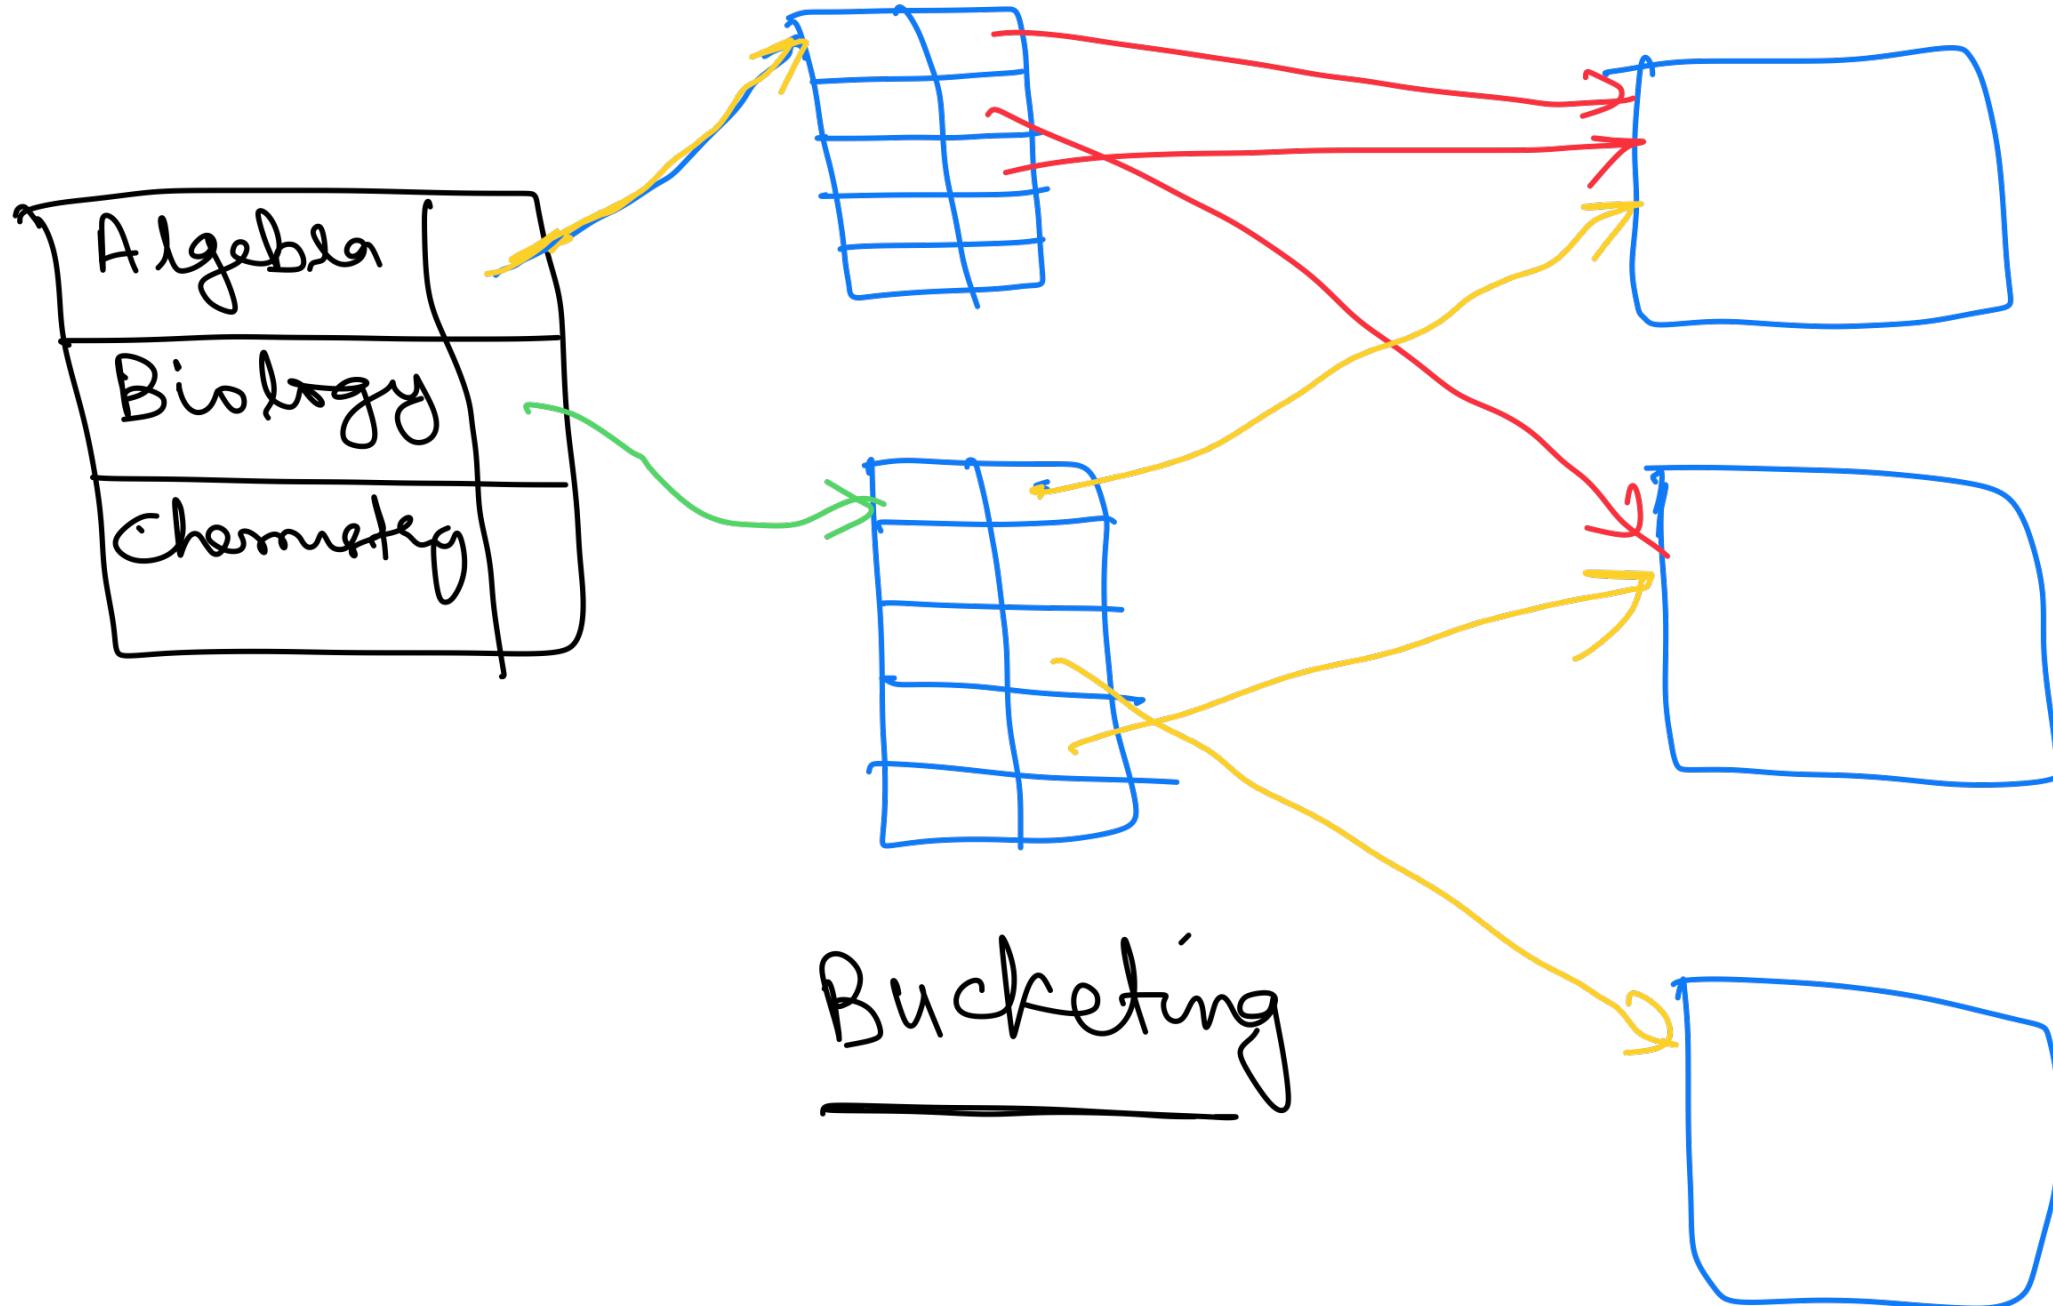
Algebra	
Biology	
Chemistry	



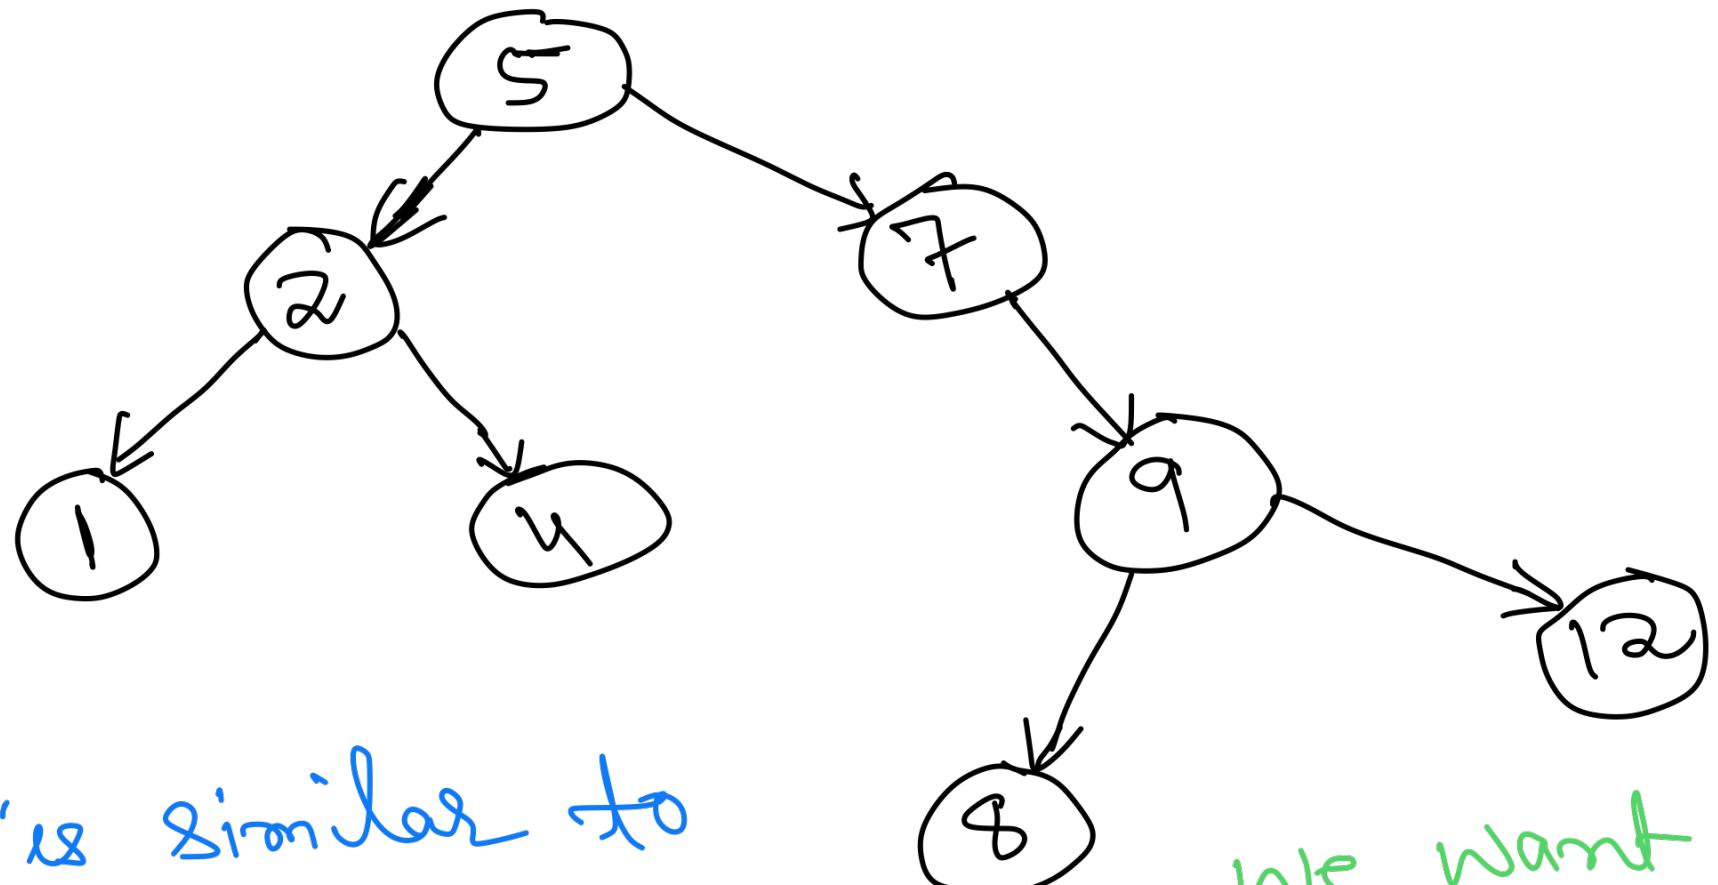
Secondary Index

---



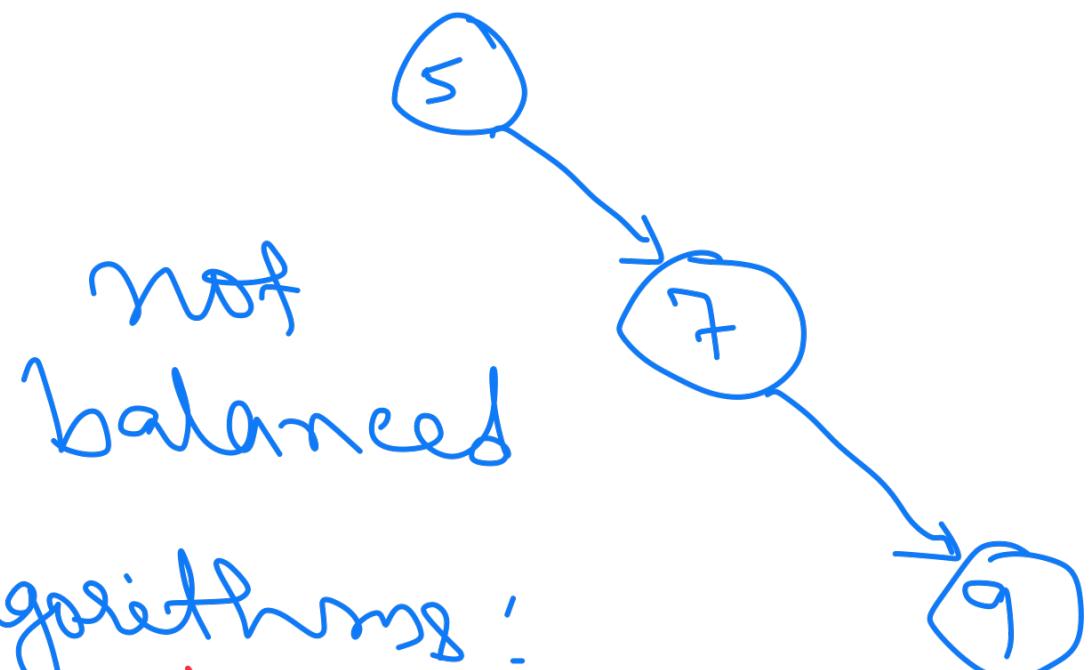


## # Binary Search Tree (BST)



It's similar to  
multi-level indexing. This BST to be  
balanced. We want  
 $n$  elements,  
 $\text{height} \leq \log_2 n$

→ BST would be balanced when for  $n$  elements, the height of this BST should not be more than  $\log_2 n$ .

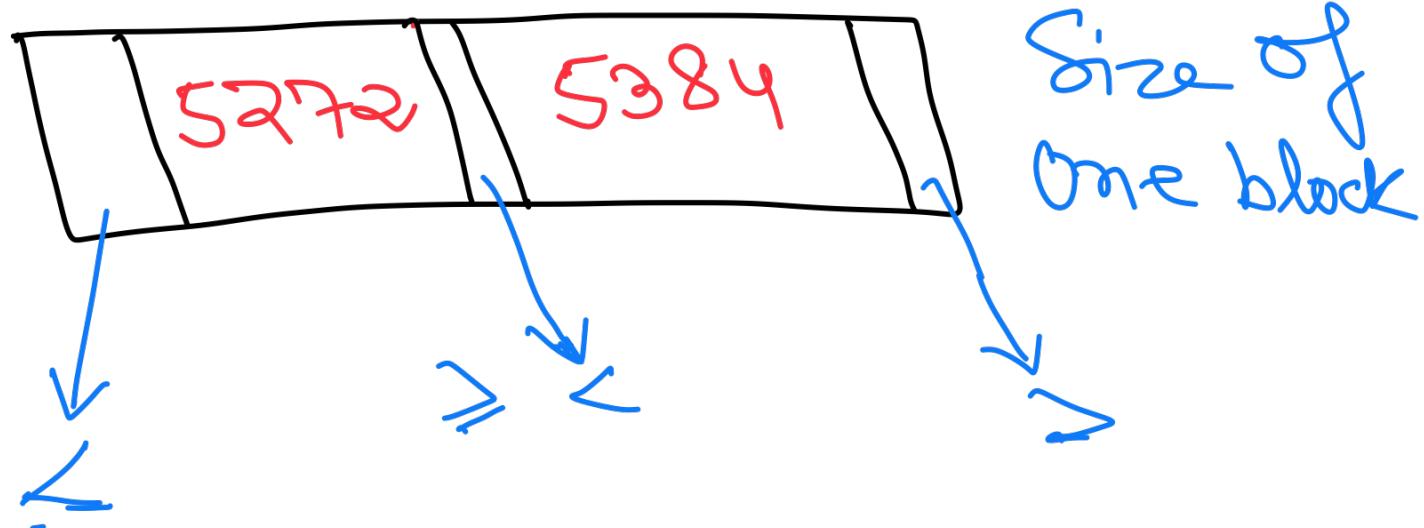


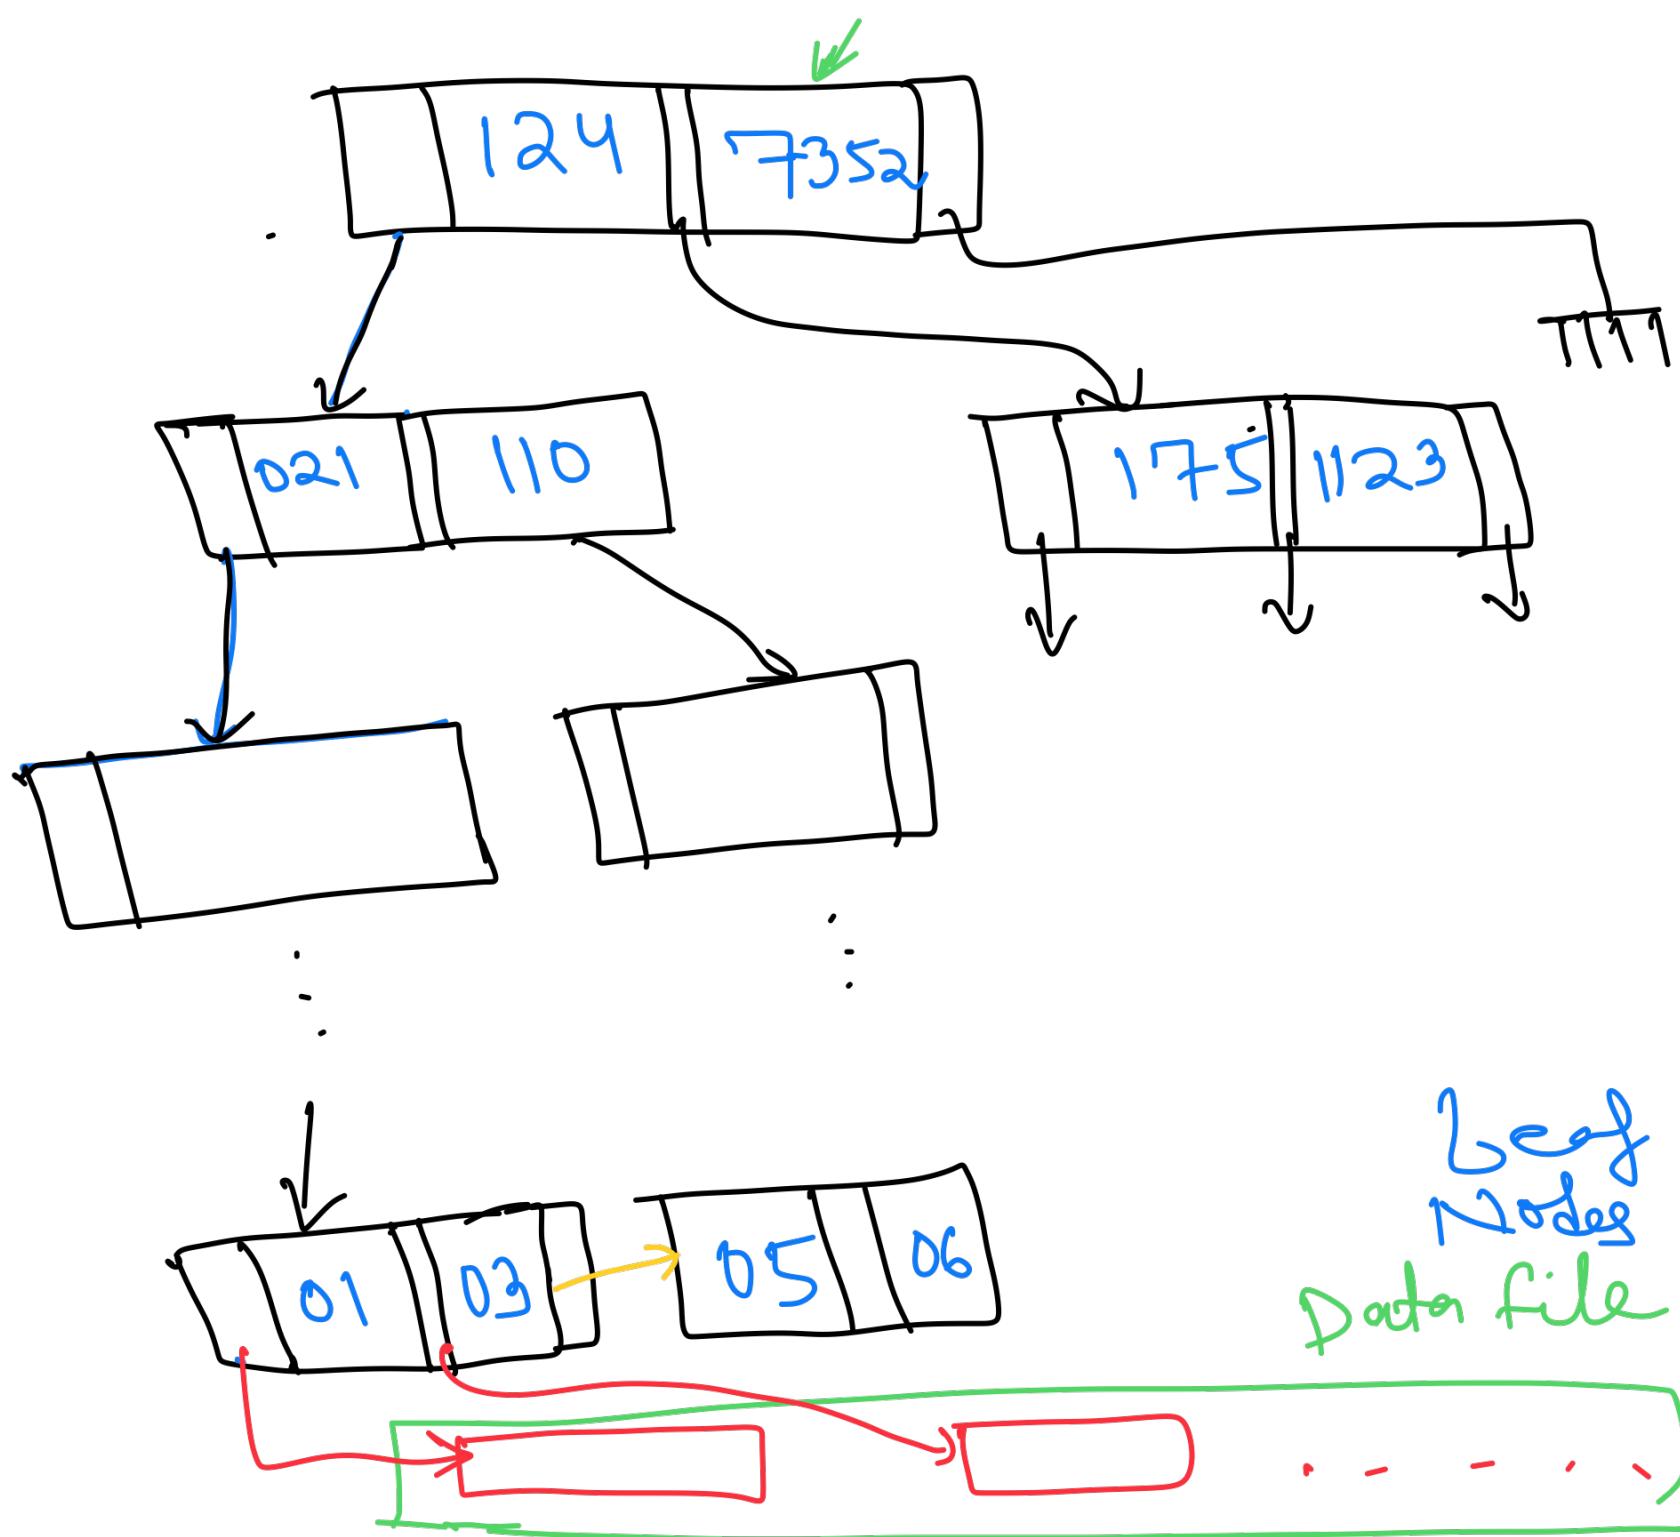
# Balancing Algorithms:  
AVL, B-Tree, B<sup>+</sup>-Tree

# # B<sup>+</sup>-Tree Index files

K Records : Each record corresponds to a key

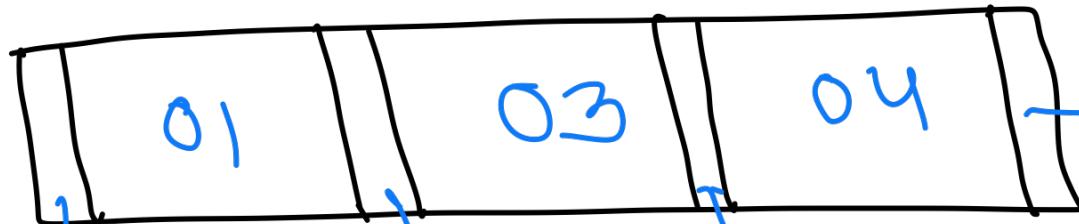
K+1 Pointers



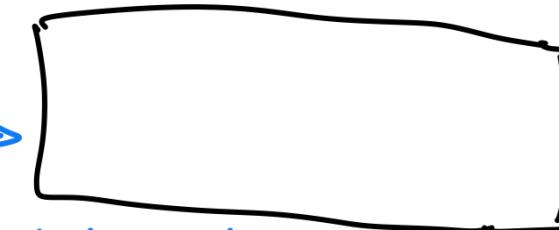


#

Leaf



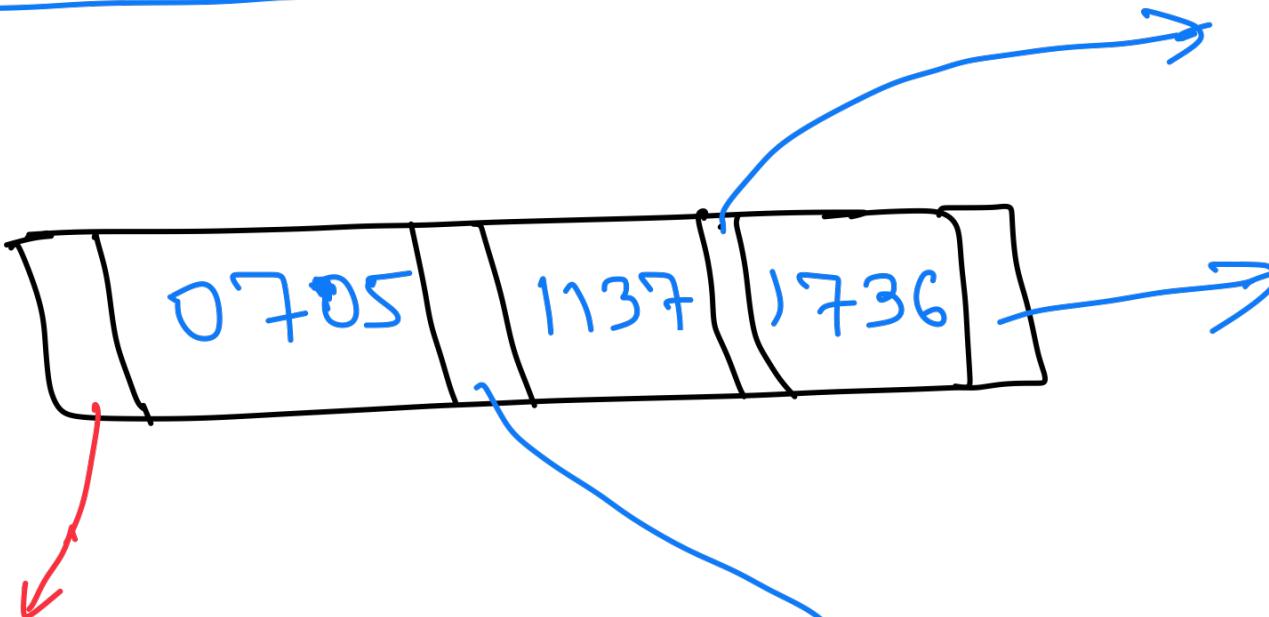
Block for  
Record with  
key 01



Next  
Leaf  
Node

#

## Non-leaf Node



Index block

(Leaf or Non-leaf)

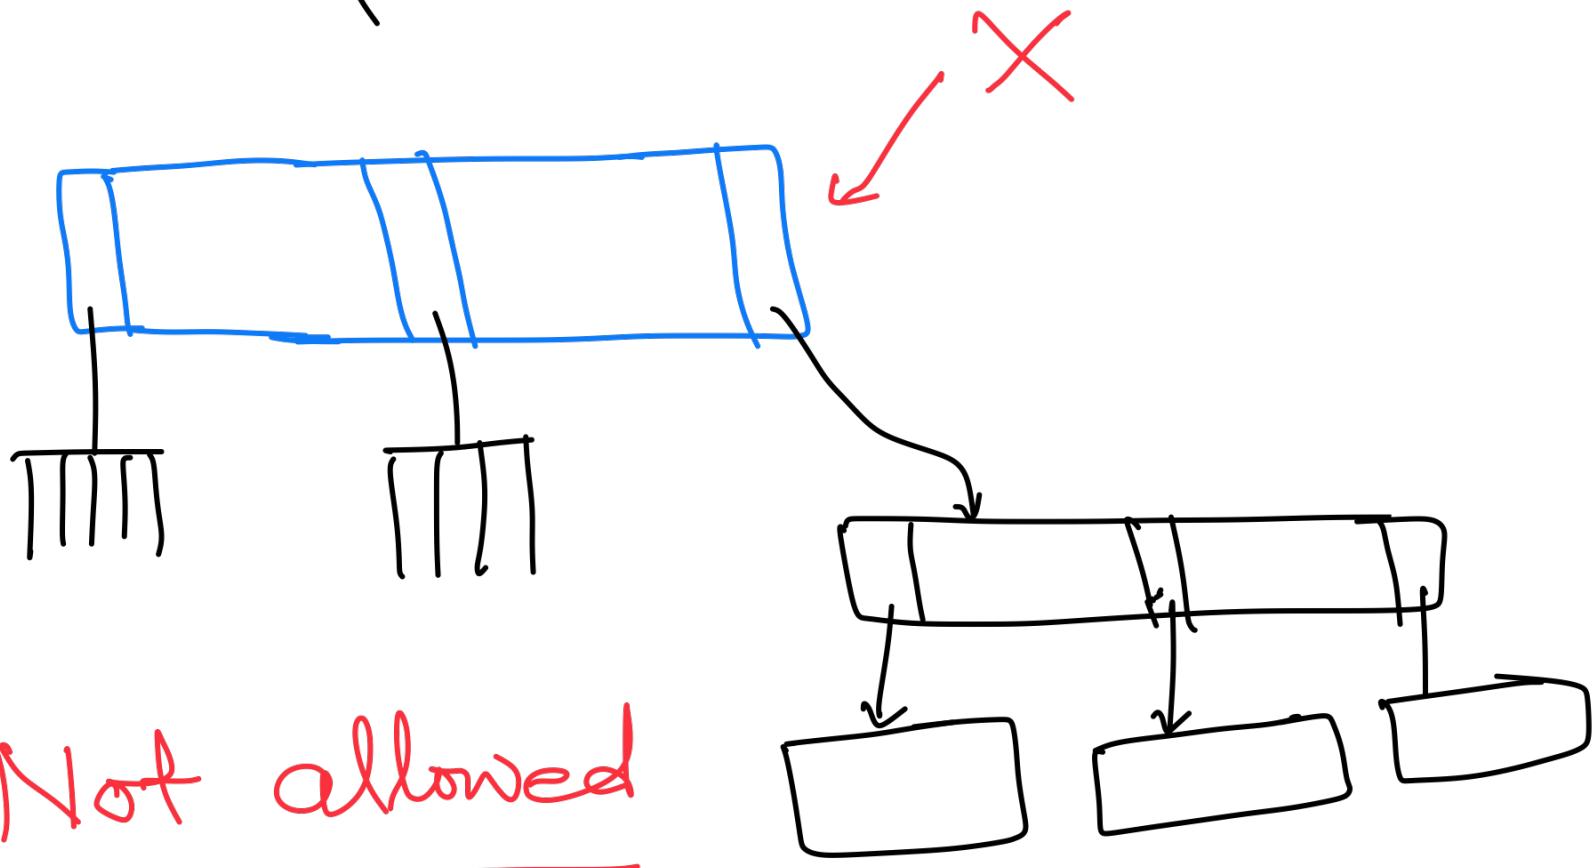
$\geq 0705$   
 $< 1137$

- for records with

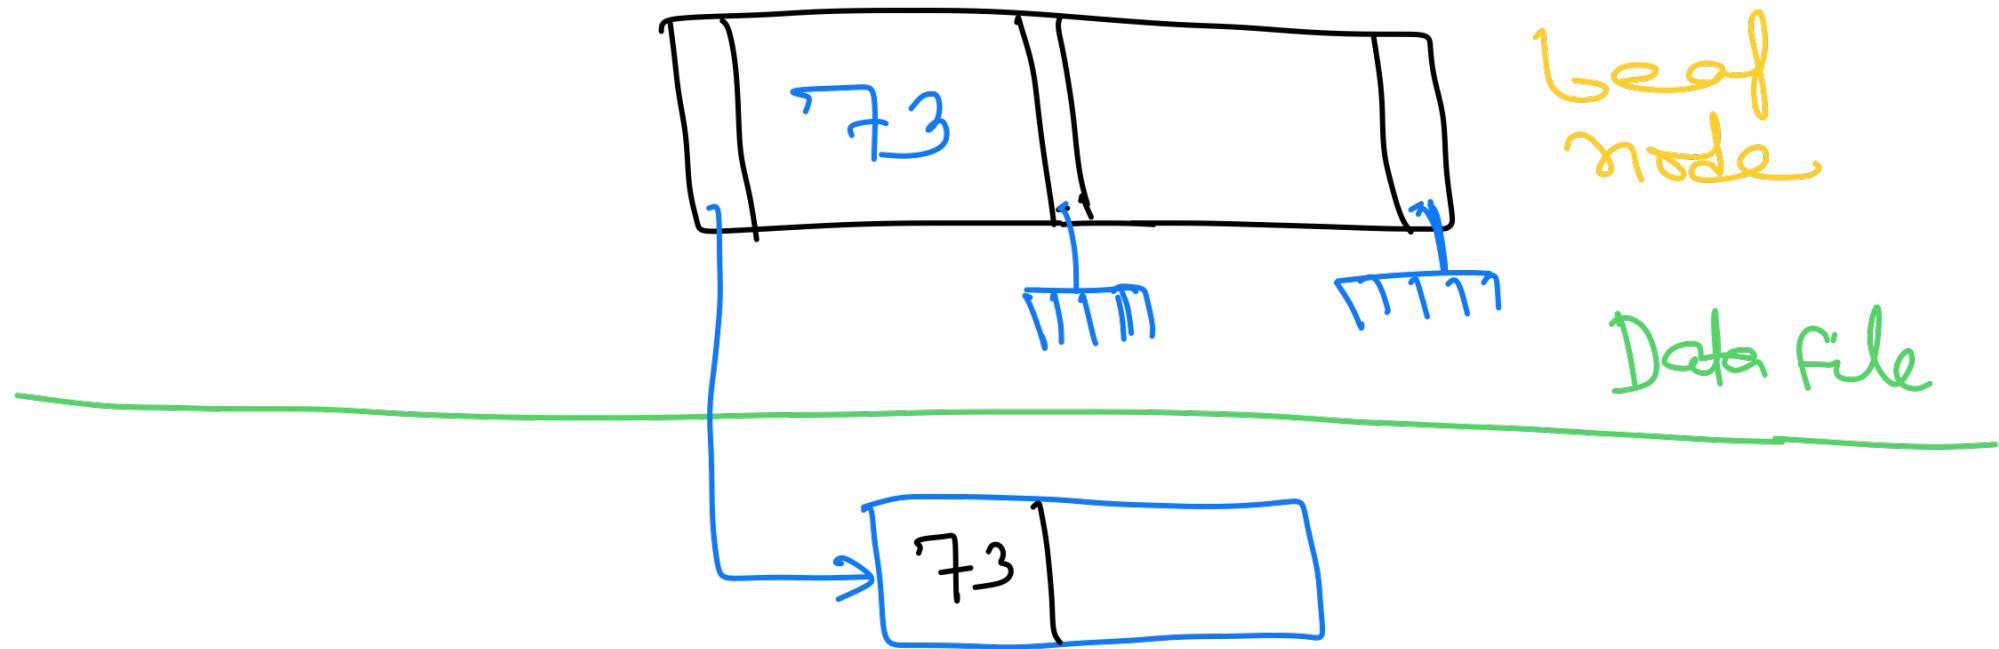
key  $< 0705$

→ All index blocks have to be at least half full.

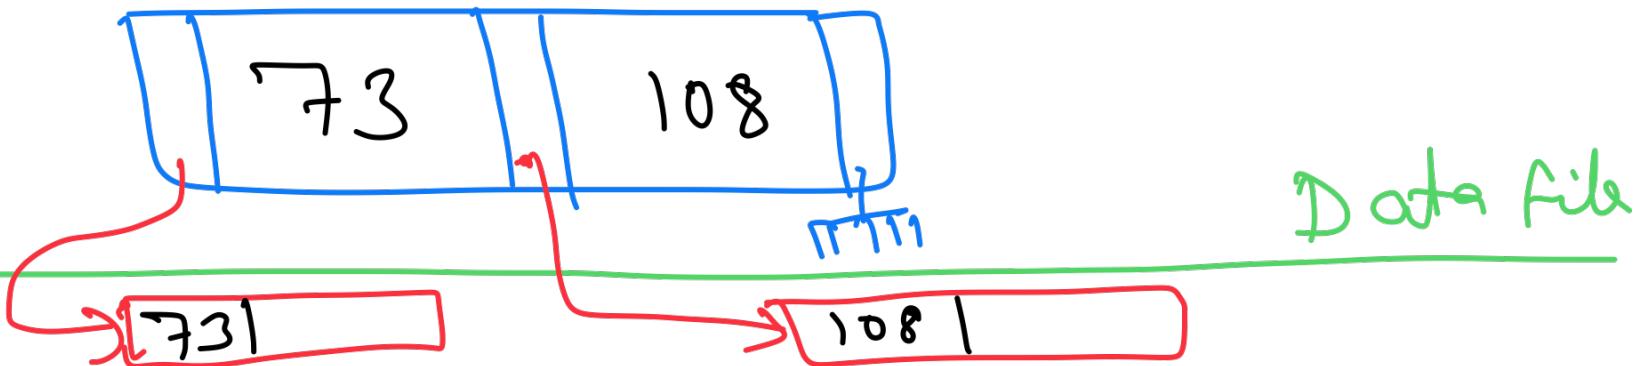
→ If  $K$  pointers are there,  
there should be atleast  $\lceil \frac{K}{2} \rceil$   
pointing to valid blocks.



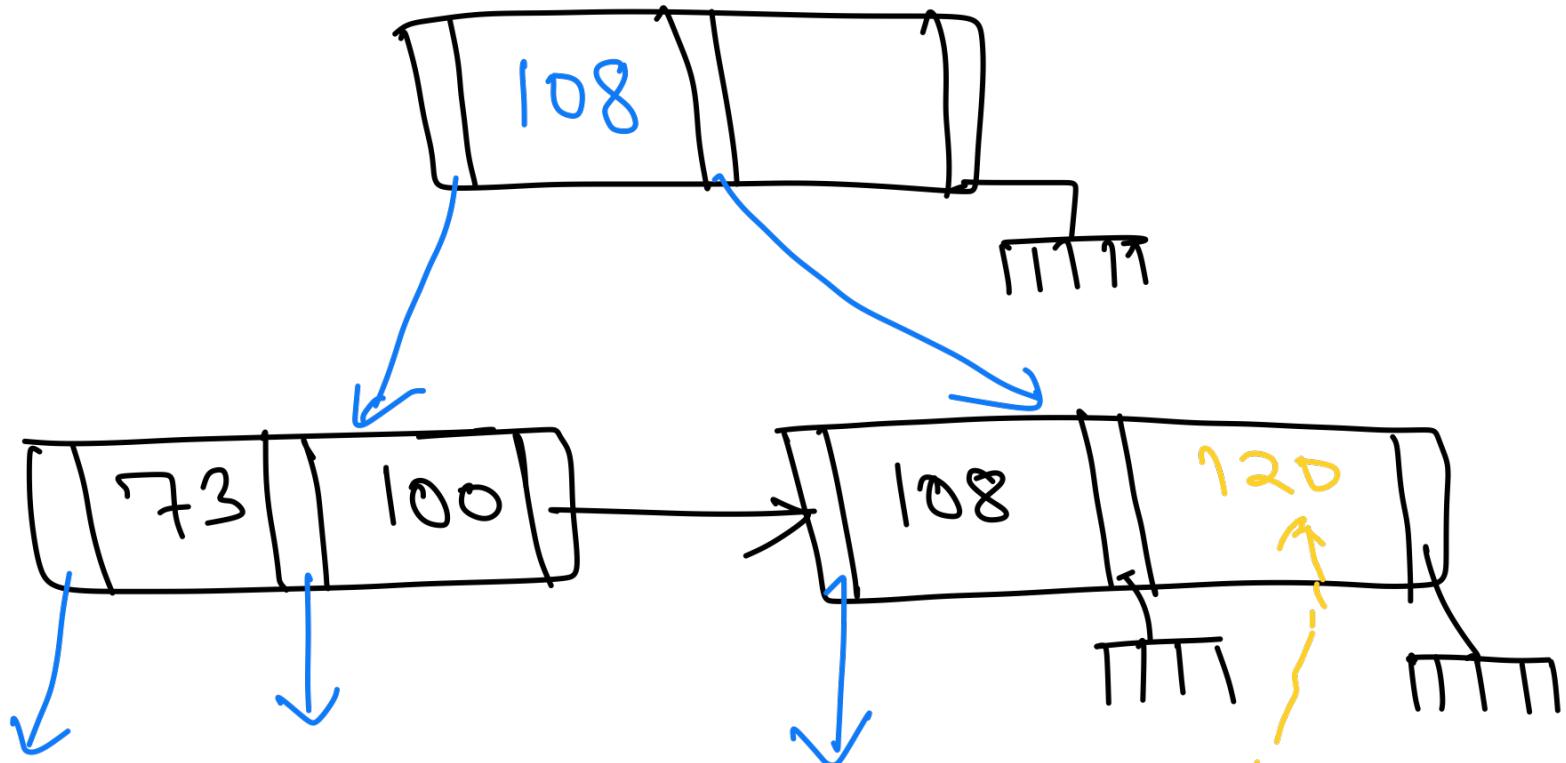
Insert 73



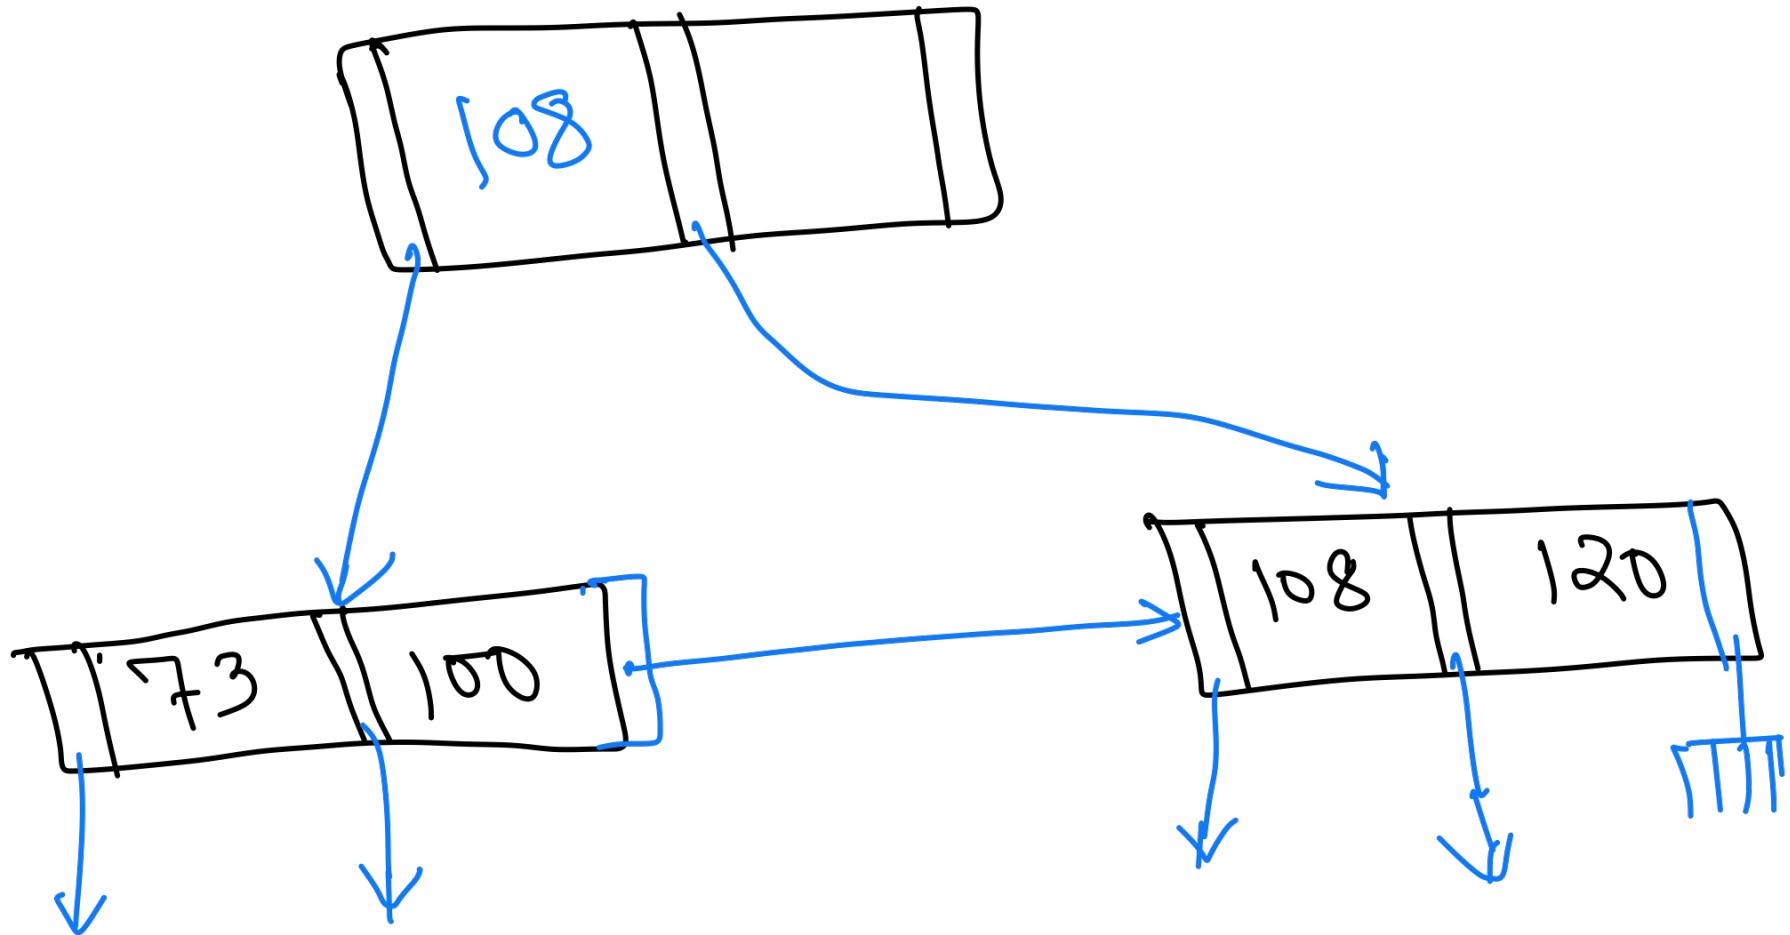
Insert 108



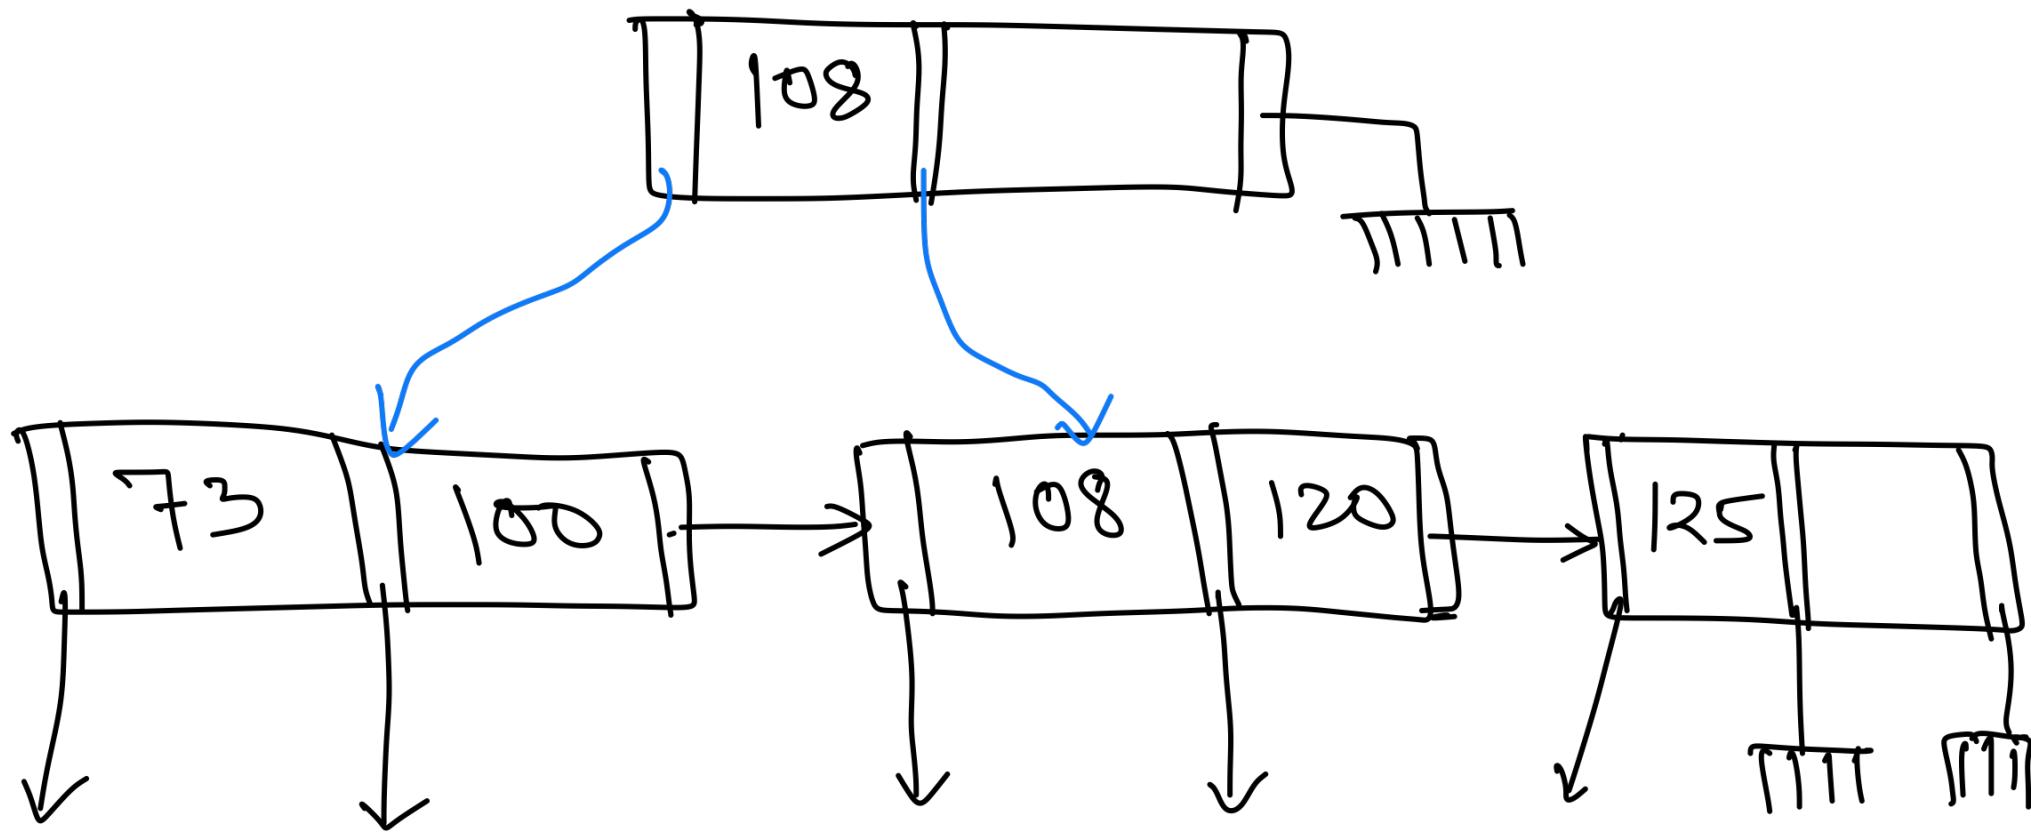
Insert 100



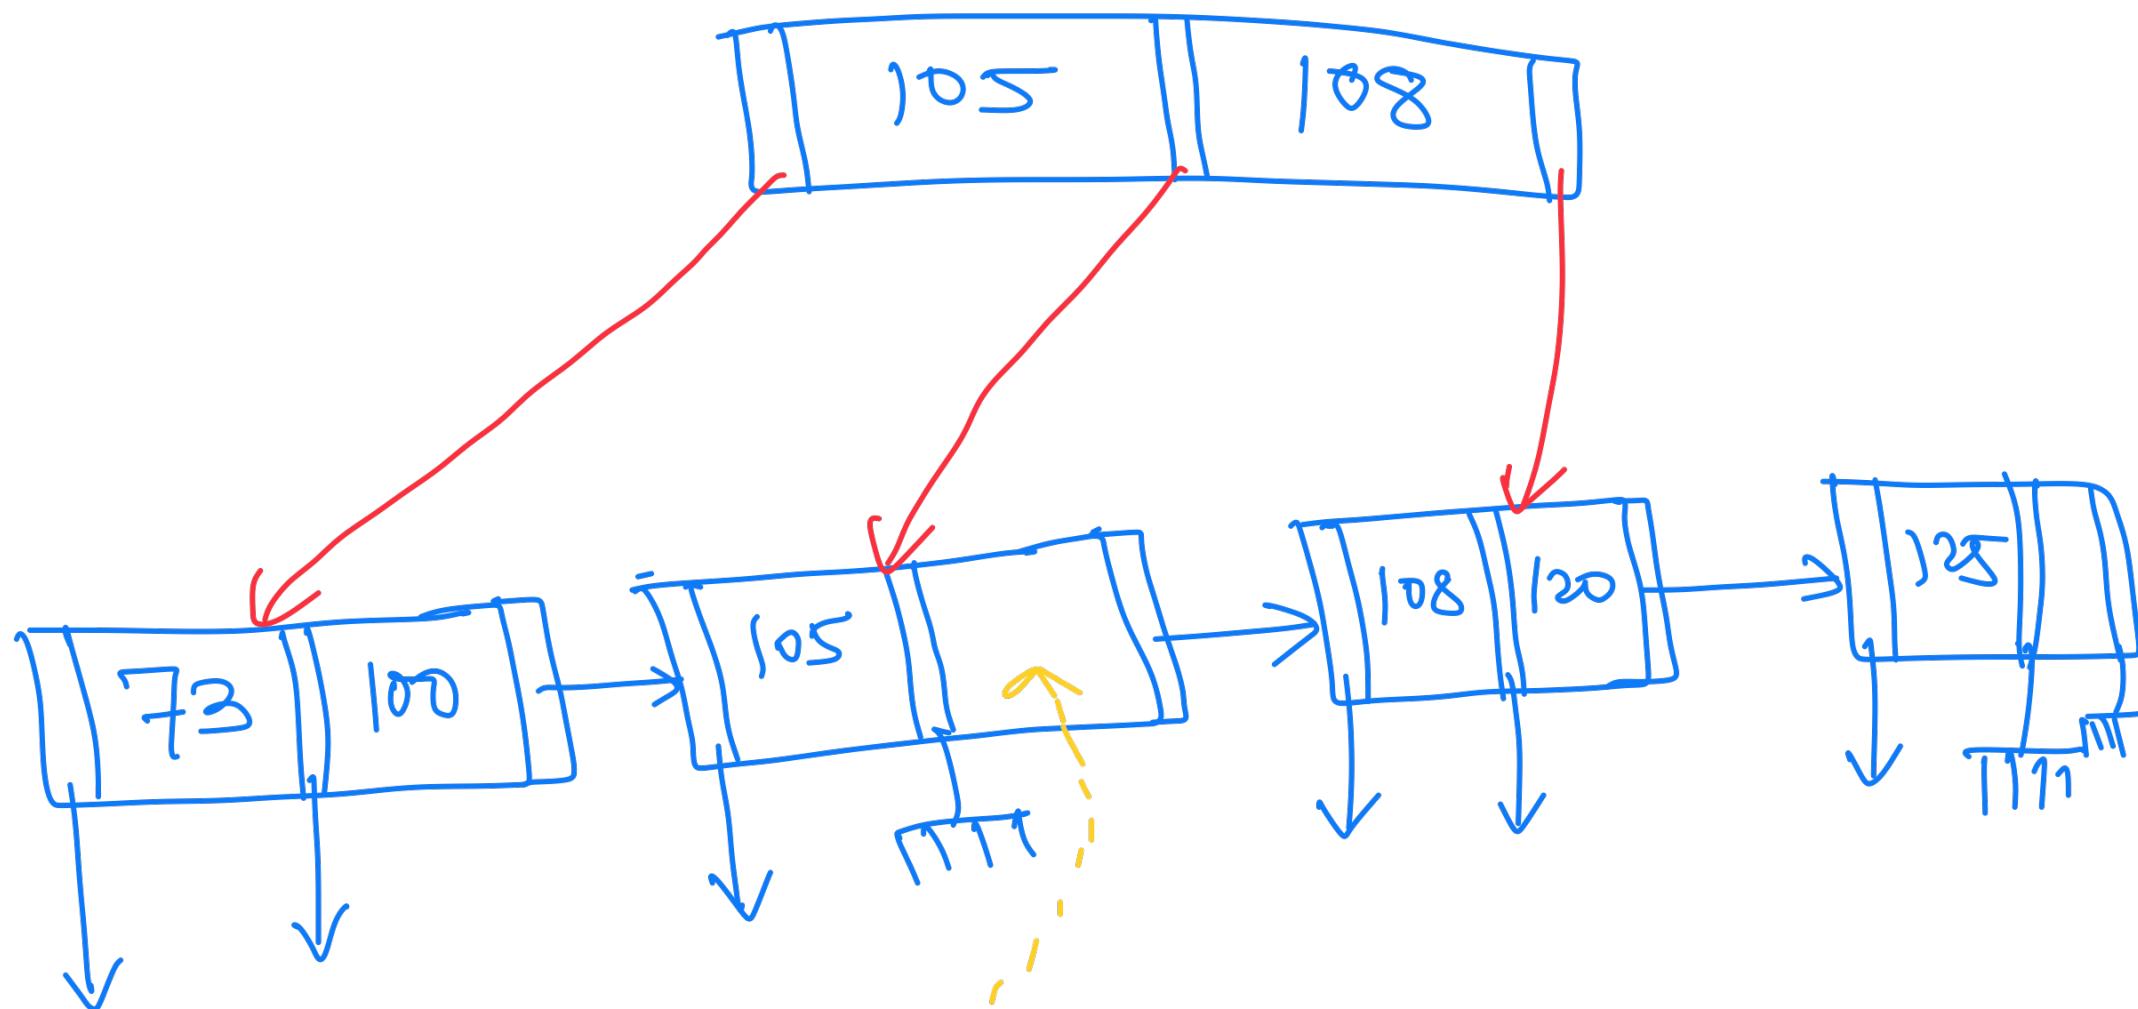
Insert 120



Insert 125



Insert 105



Insert 106

Insert 107

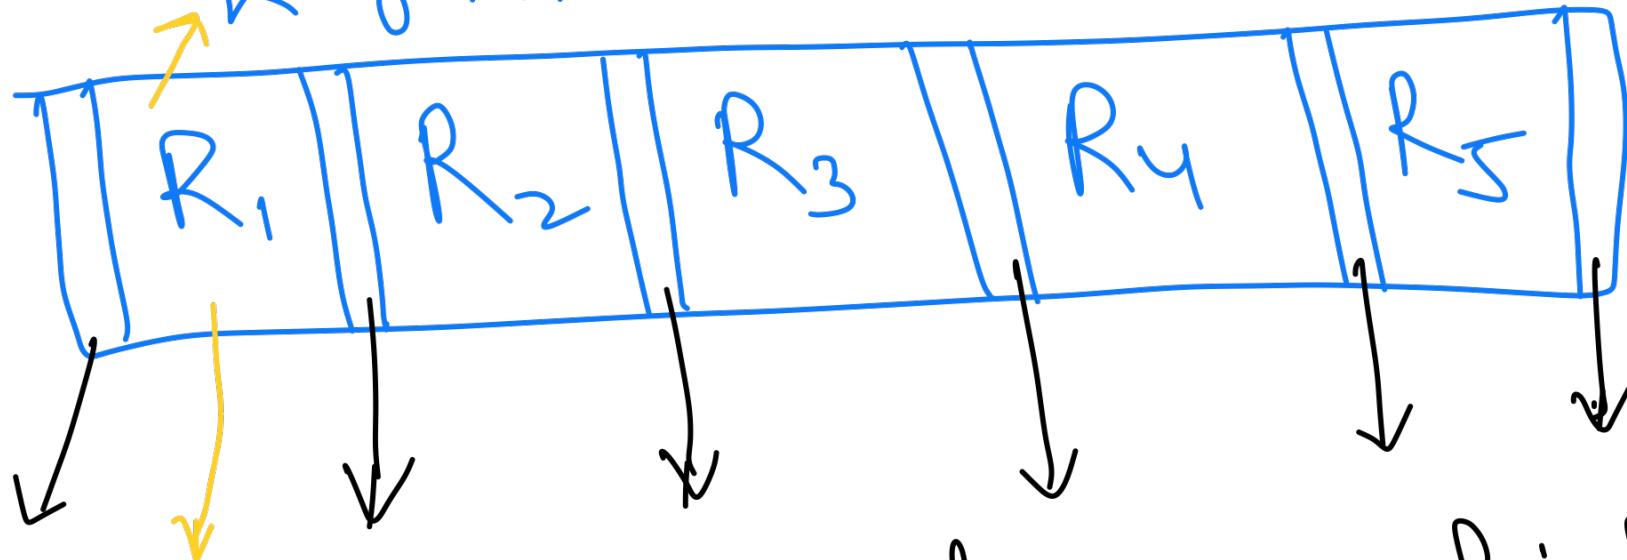
$$\log_3 n \leq \text{height} \leq \log_2 n$$

## B<sup>+</sup>-Tree

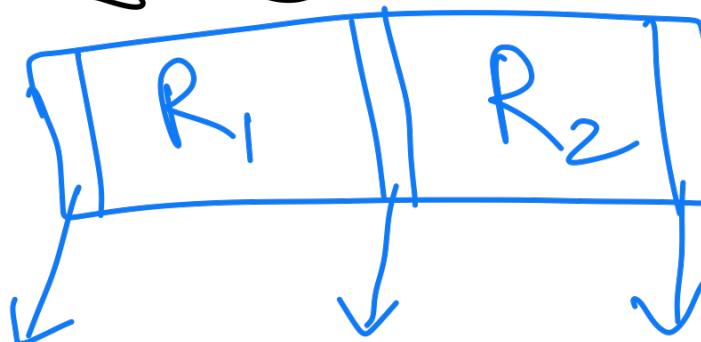
- Balanced Search Tree
- Insert , Delete , Find  $O(\log n)$
- Leaf Node and Internal Node
- All leaf nodes are at the same depth on a B<sup>+</sup>-Tree
- Each node is the size of a block

- Each node (internal node) is always at least half full.

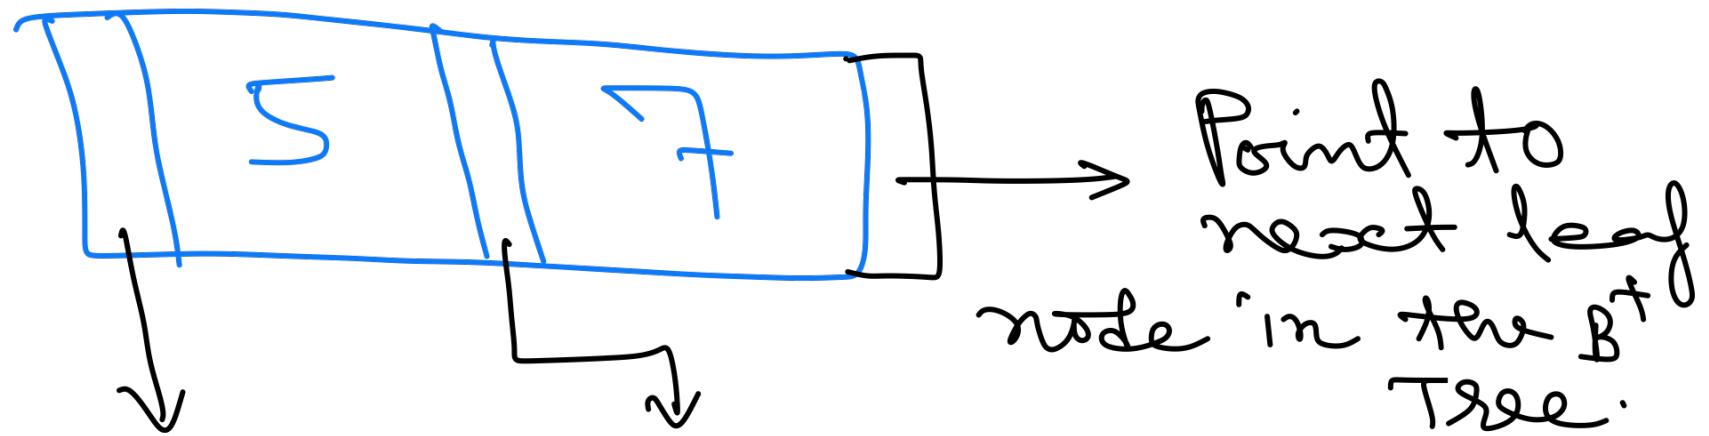
Key Attribute



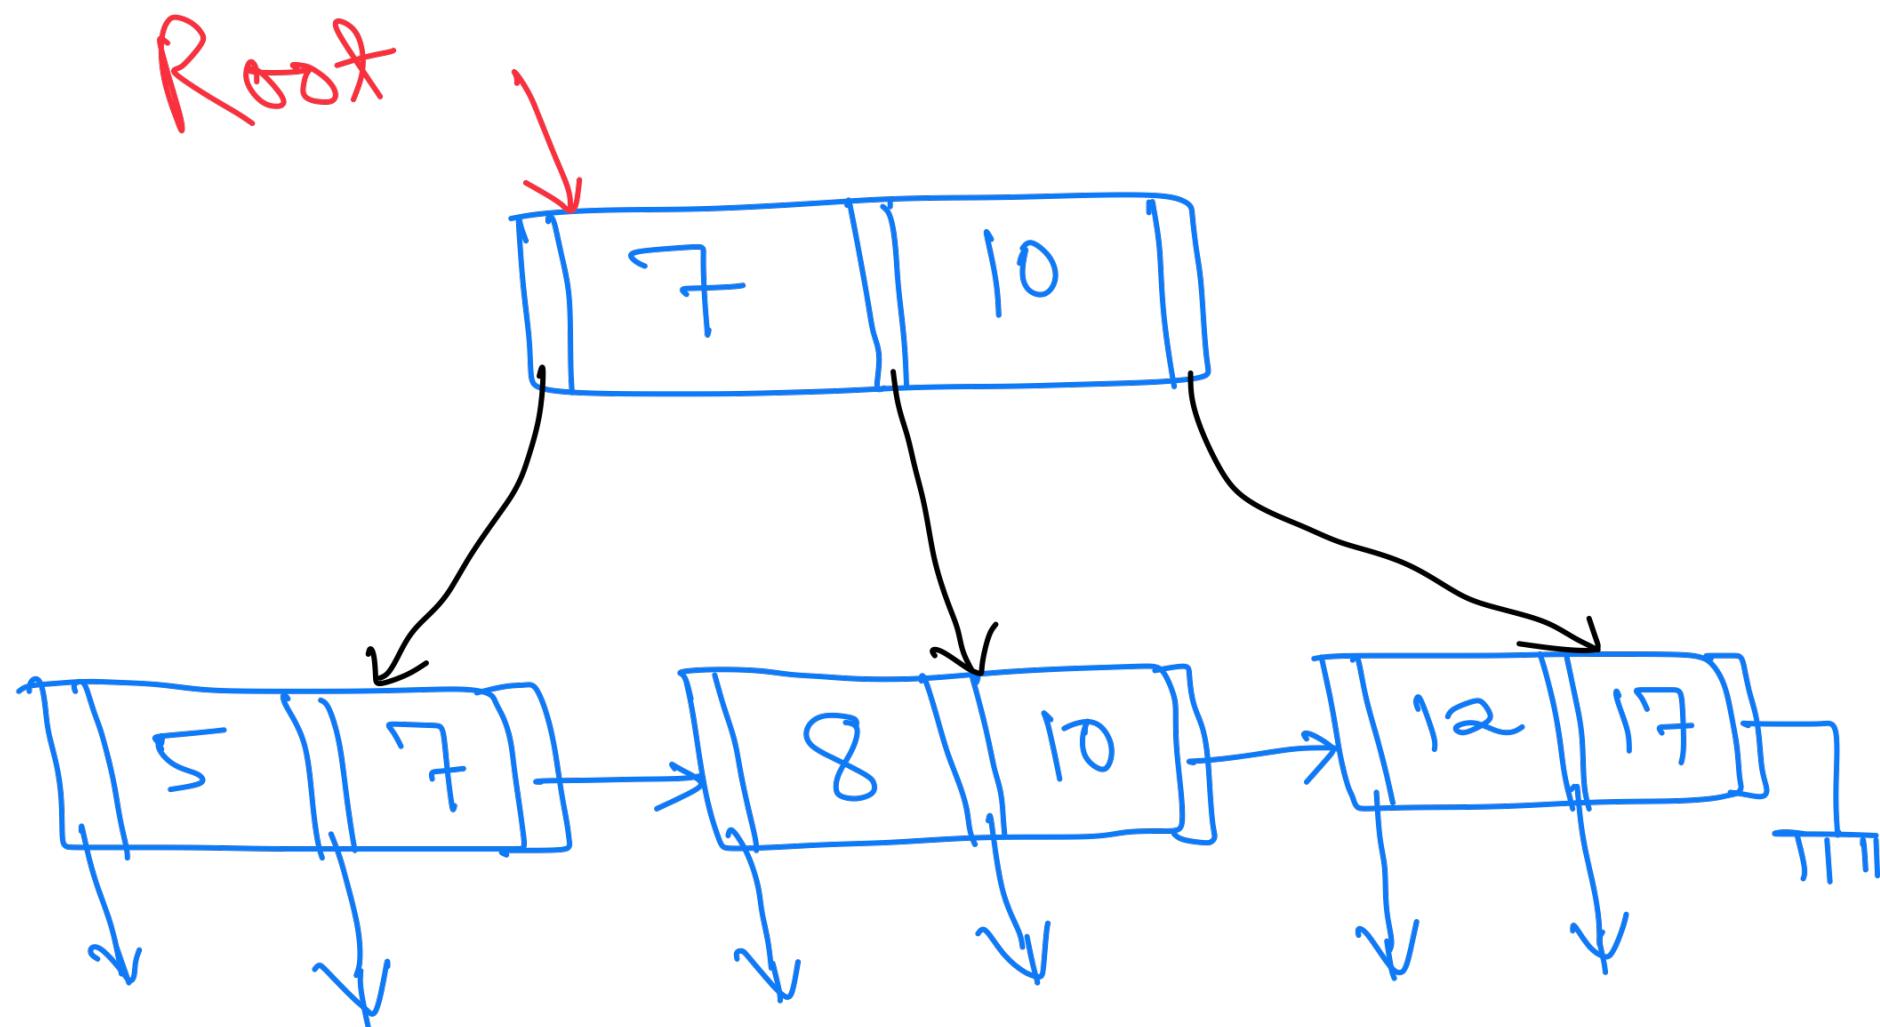
The attribute set on which indexing is done.



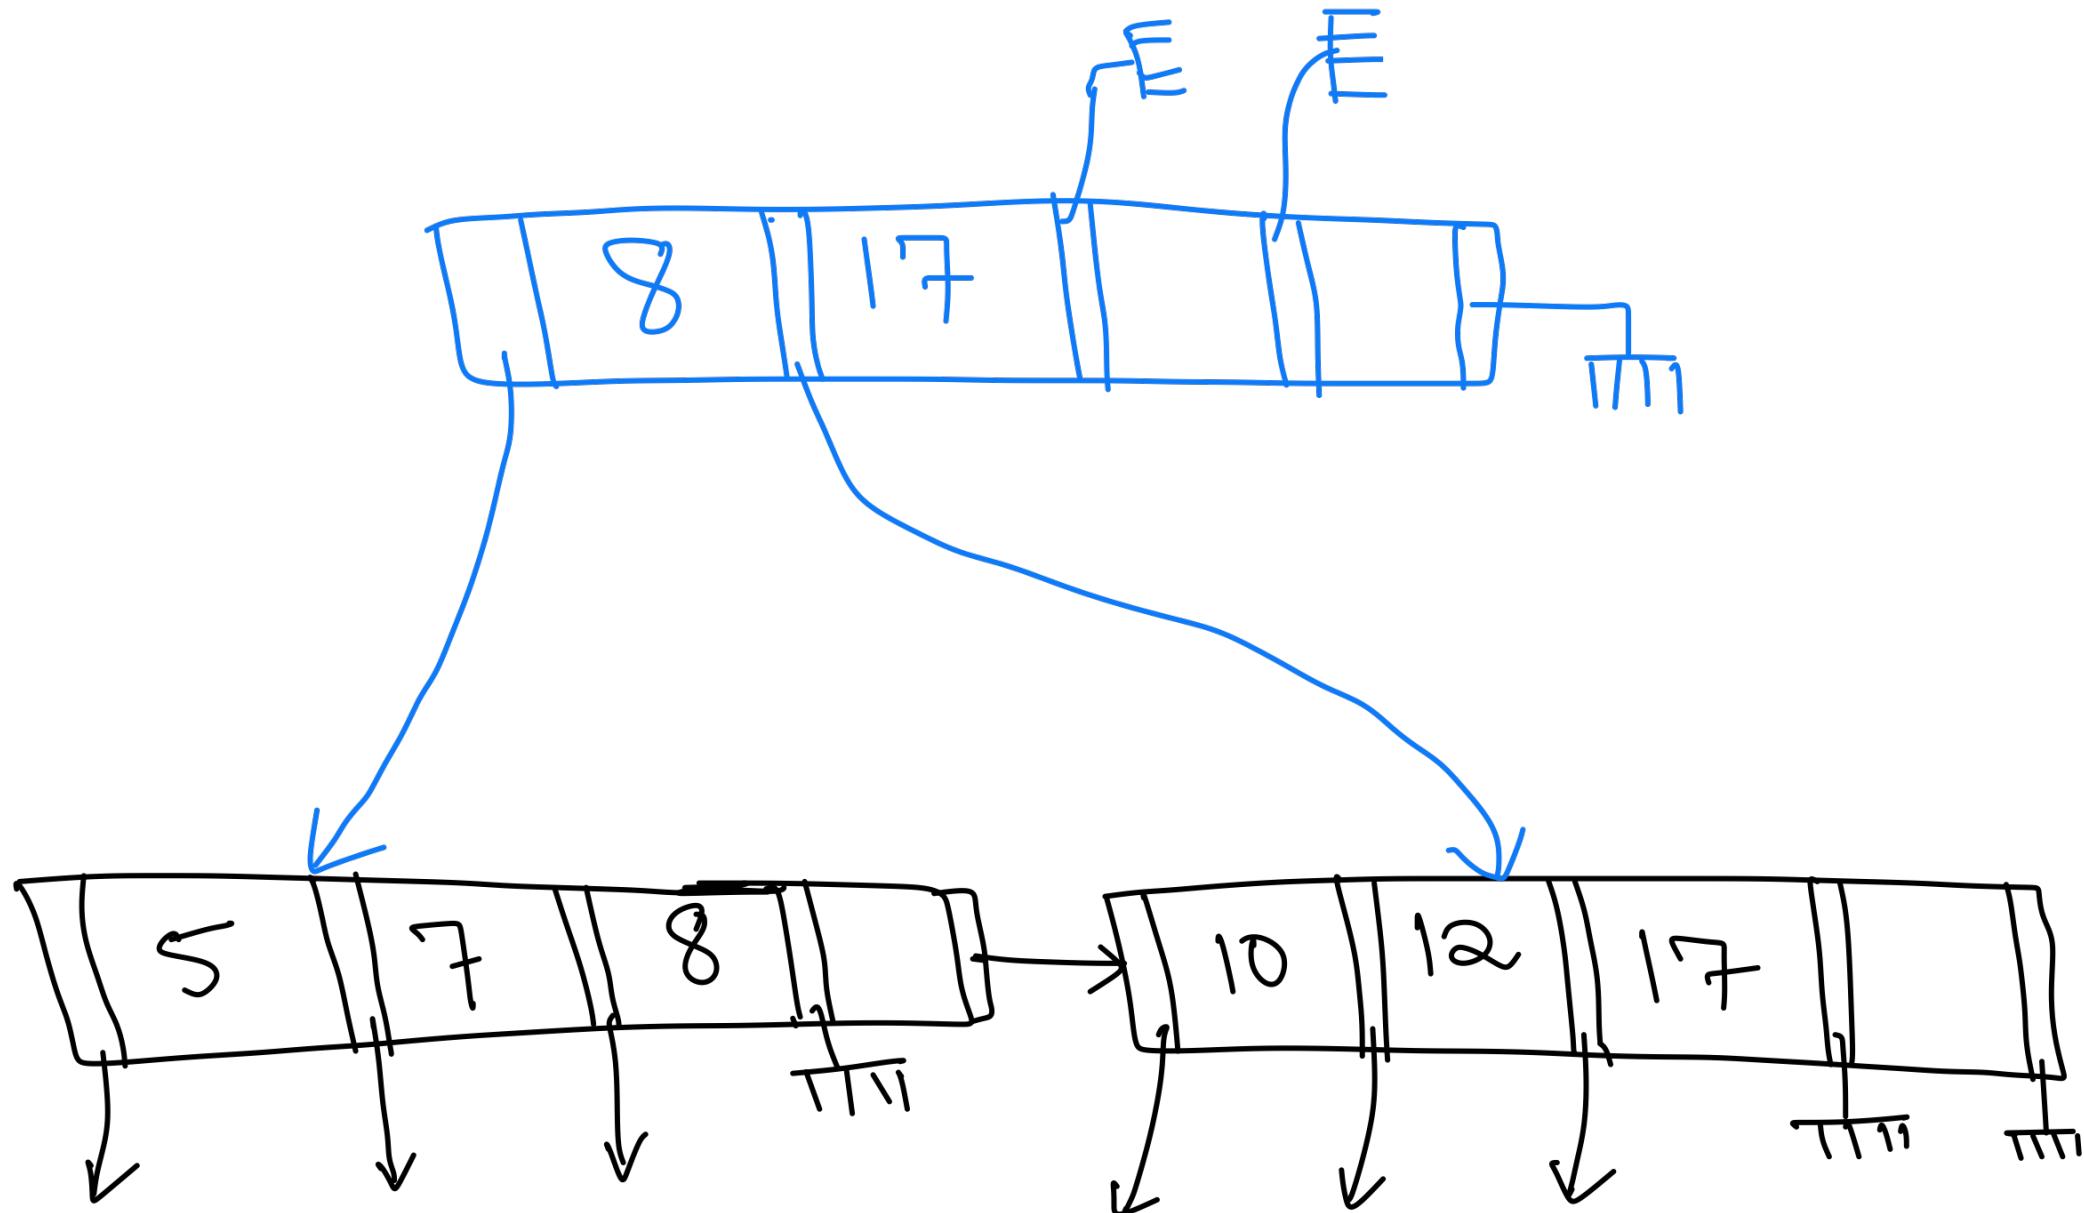
# Leaf Node



Point to the block where the record corresponding to the attribute(s) Value 5 exists in the main/Base file.



$\lfloor k, \lceil k/2 \rceil \rfloor$  must be non-null



Not a valid B<sup>+</sup>-Tree

Root

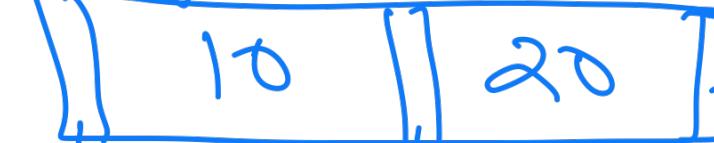


Insert (10) :-

Root → Point to block where the record exists.

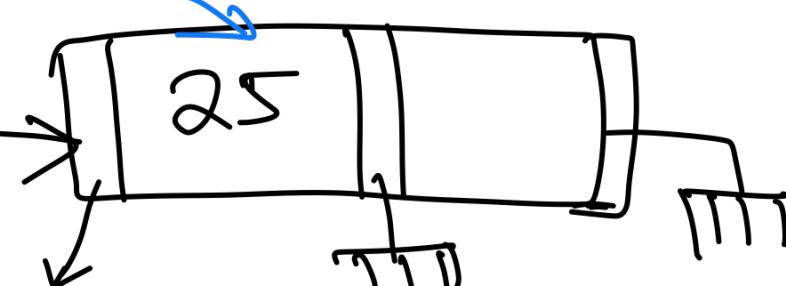
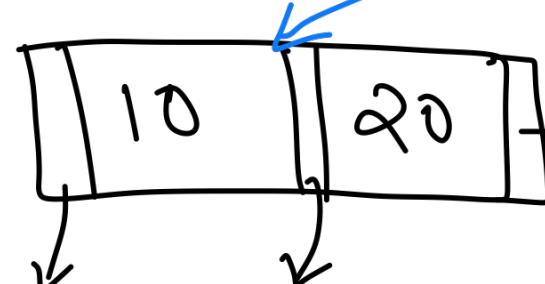
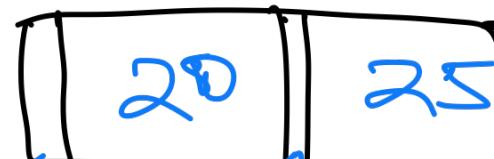
Insert (20) :

Root

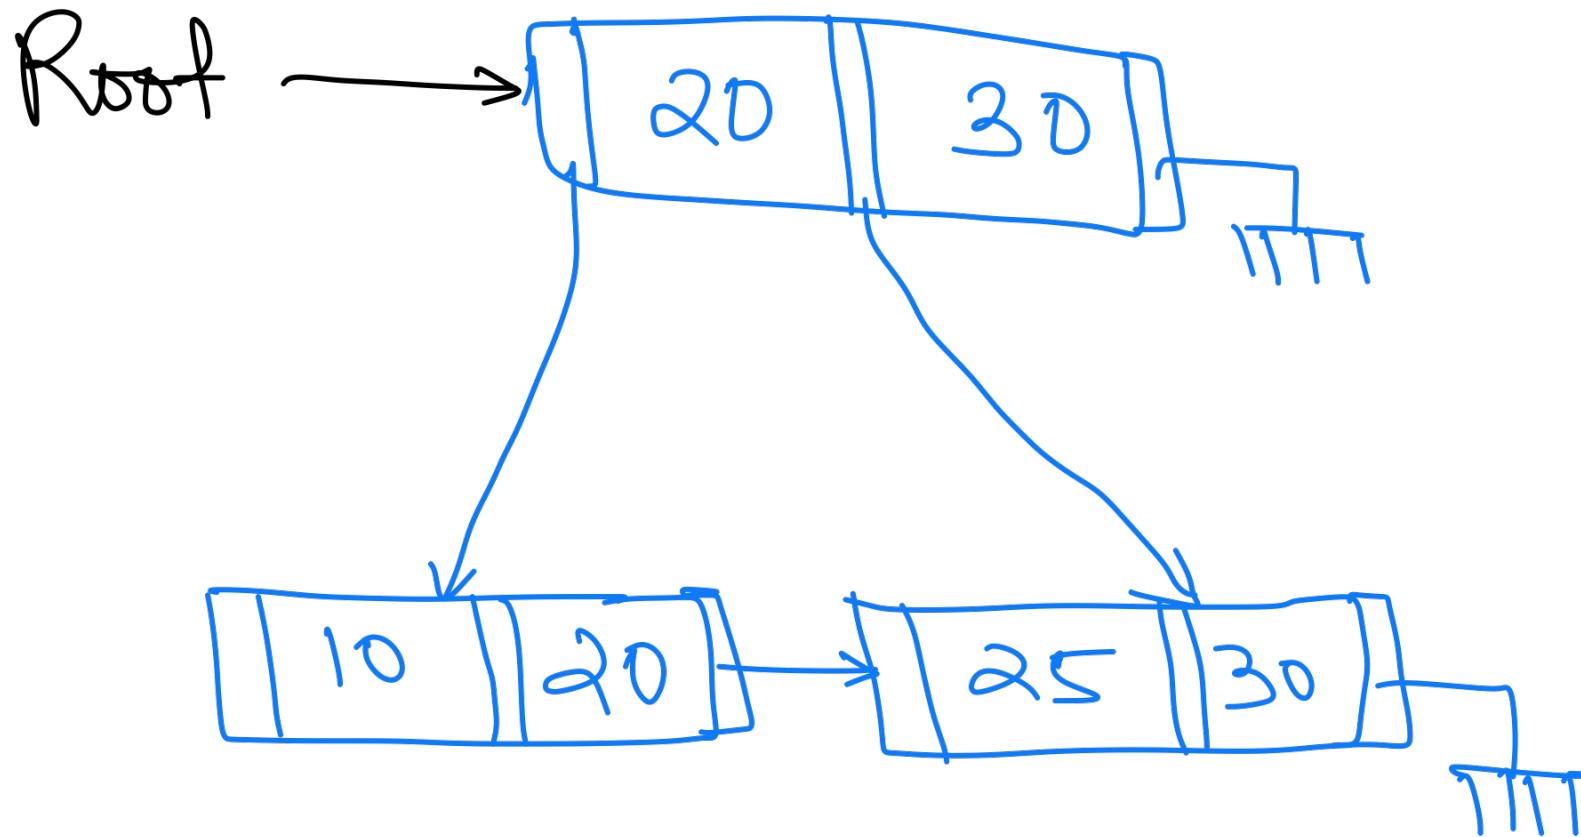


Insert (25)

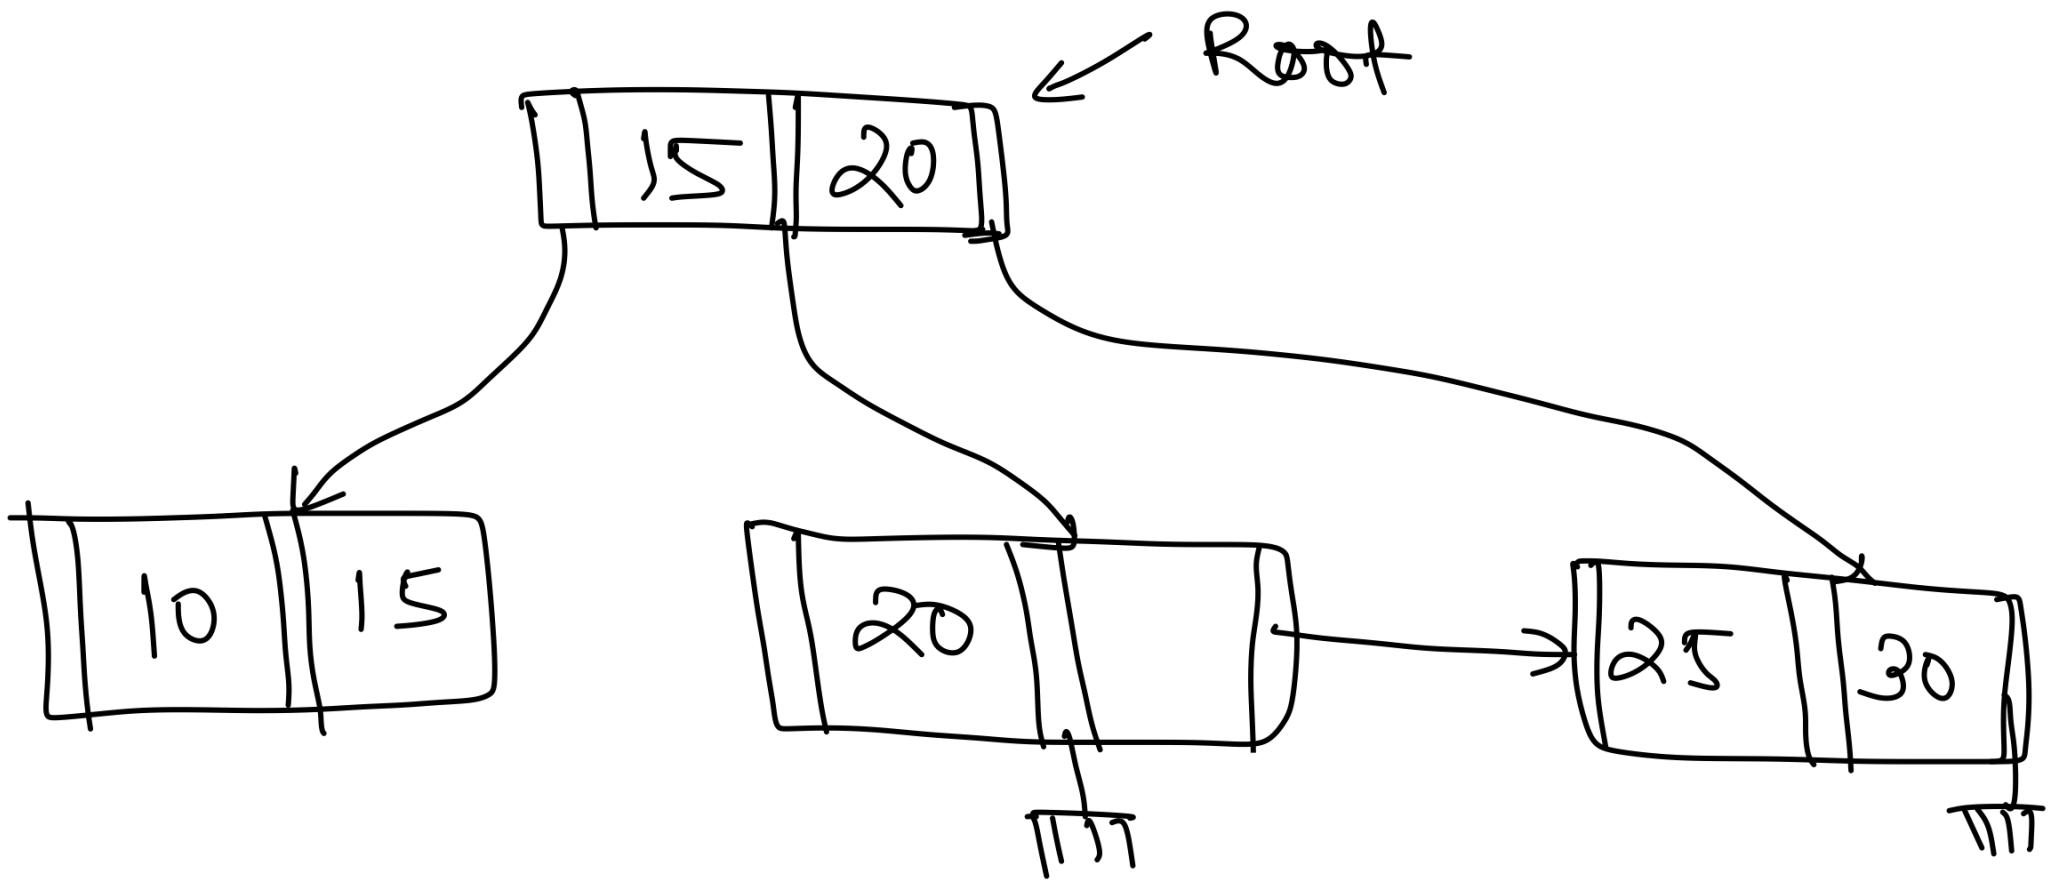
Root



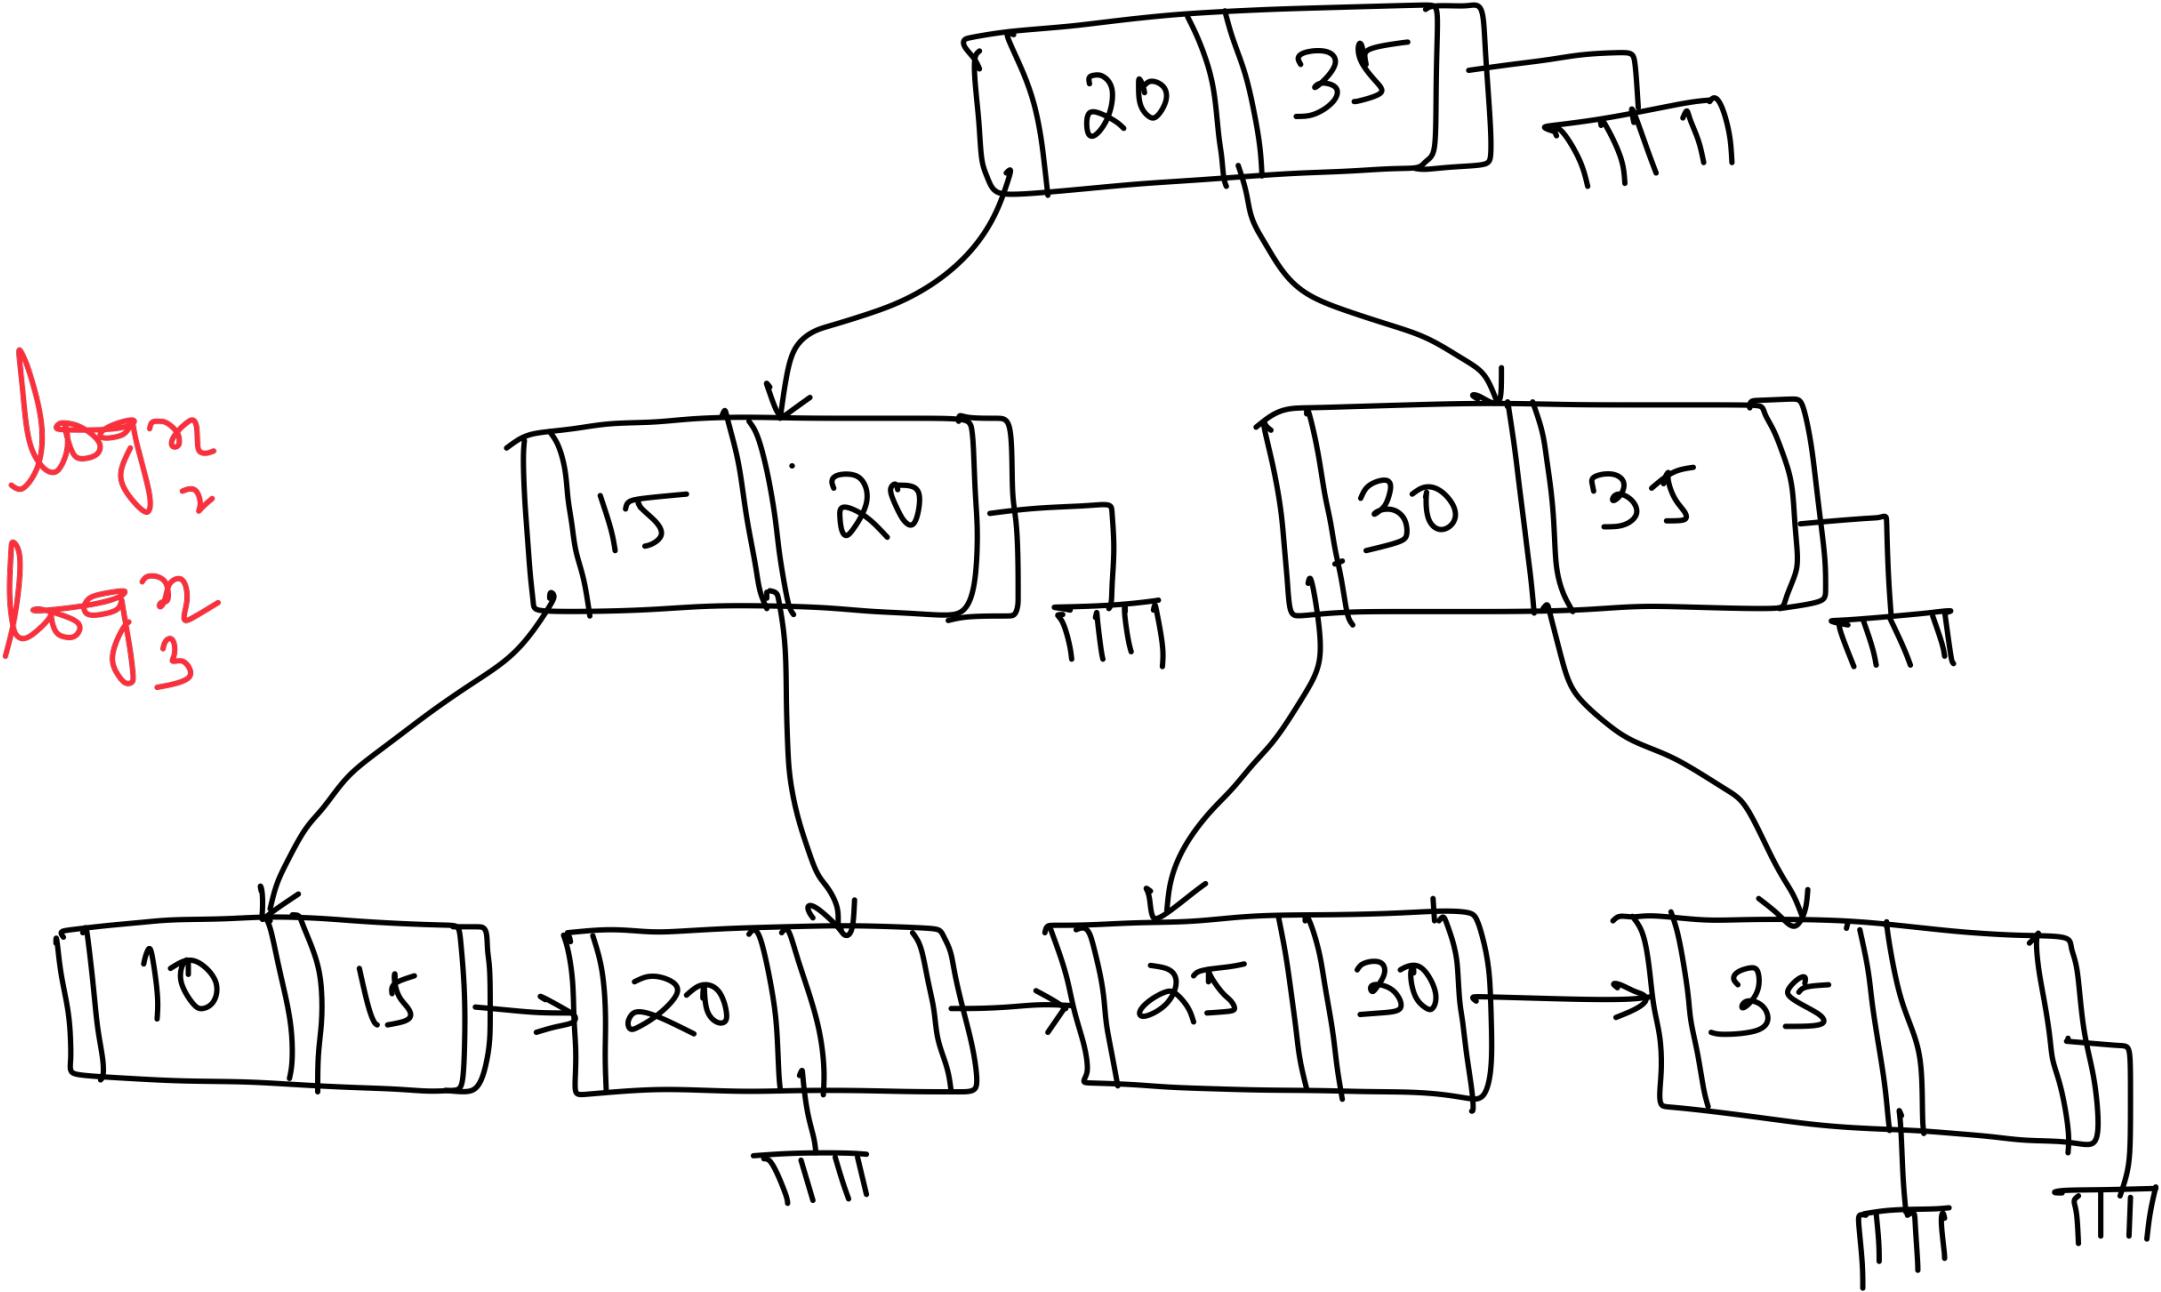
Insert (30)



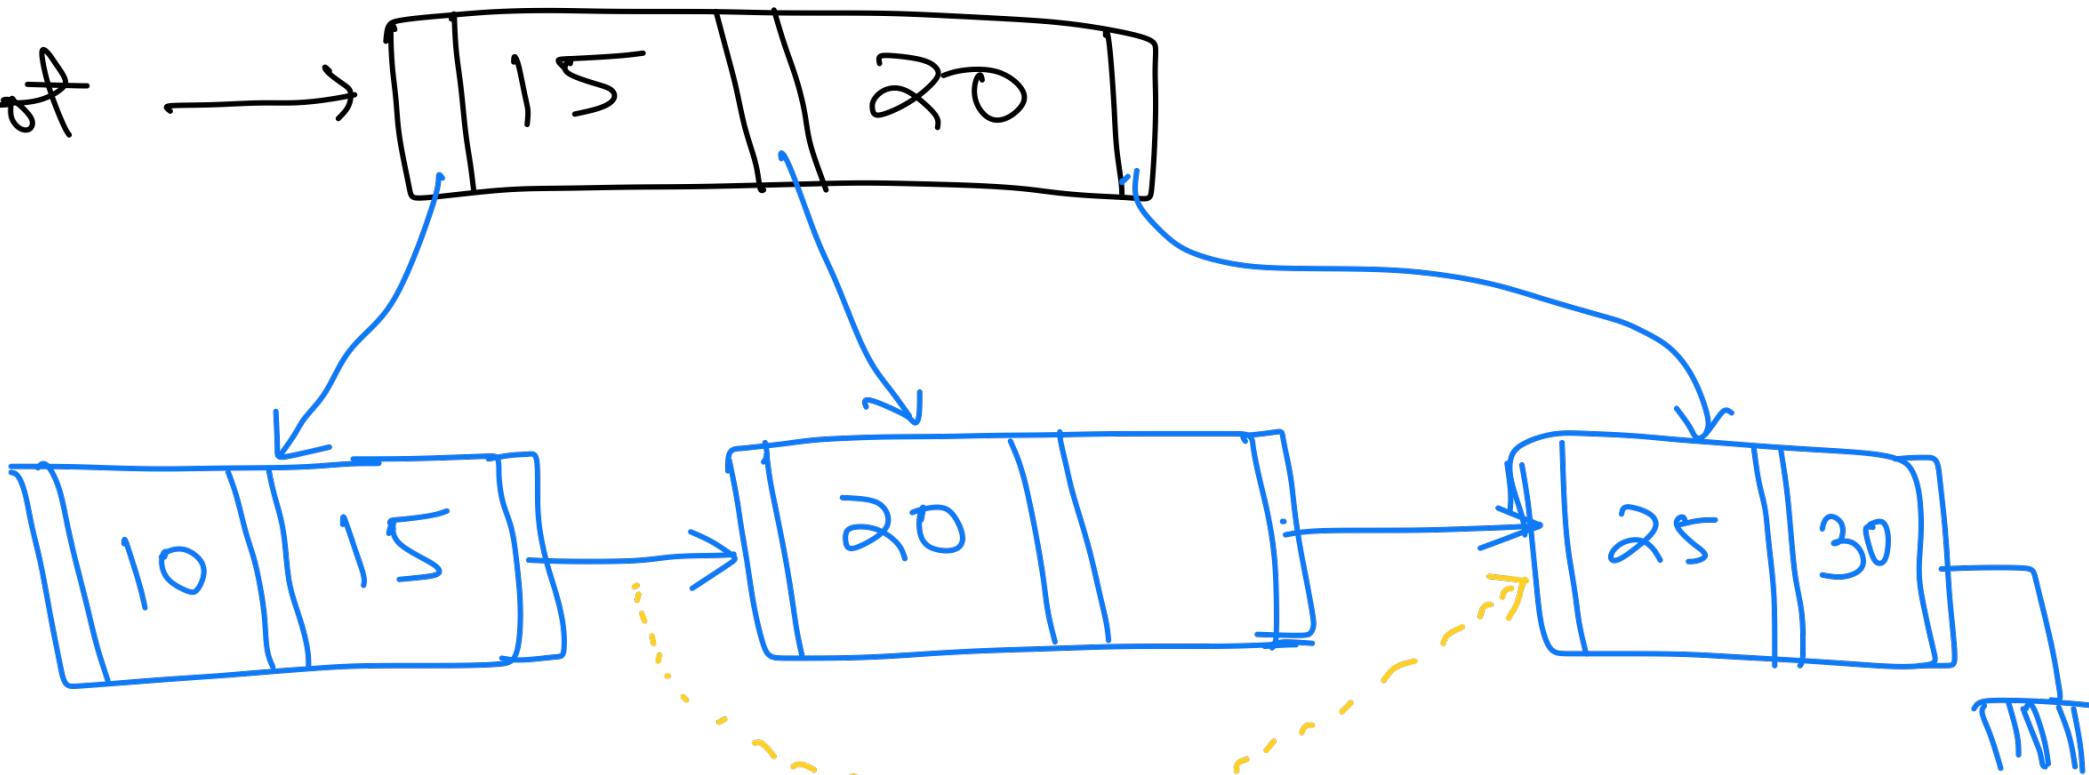
Insert (15)



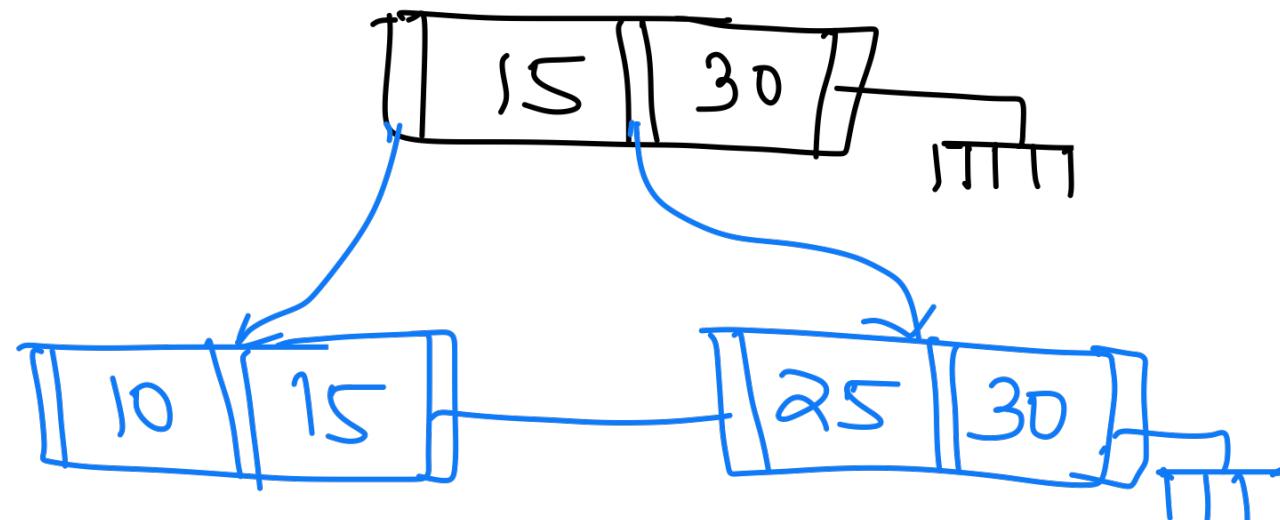
Insert 35

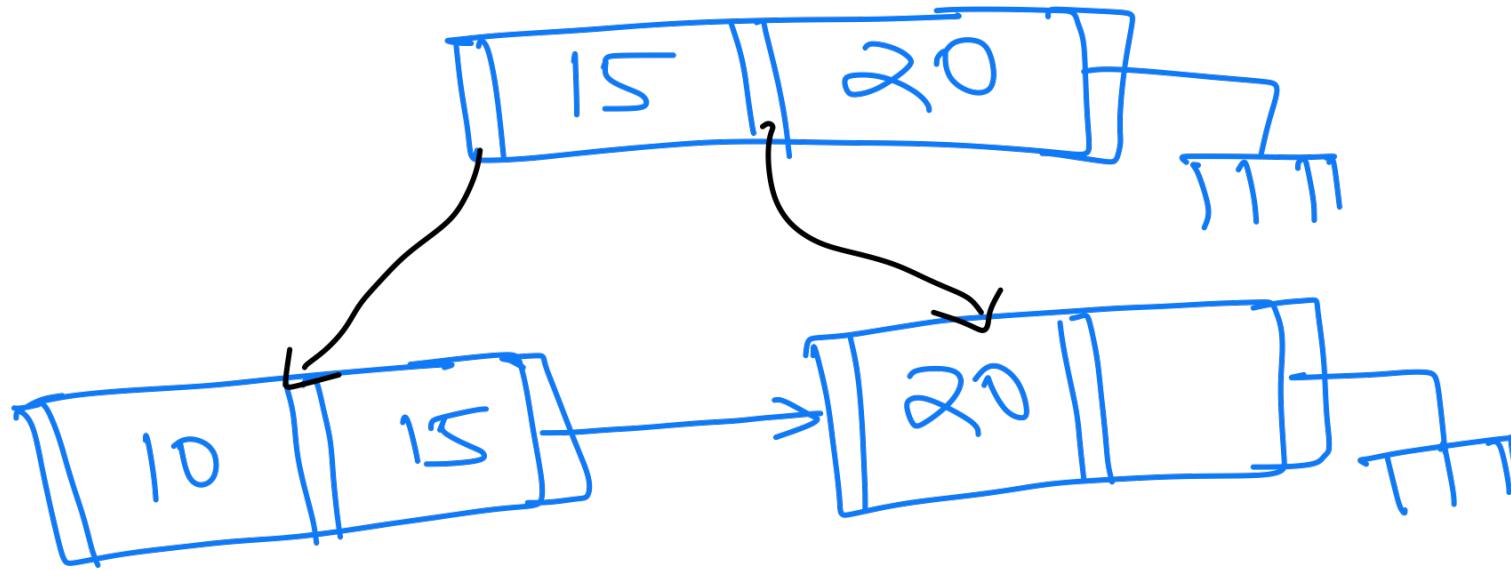


Root

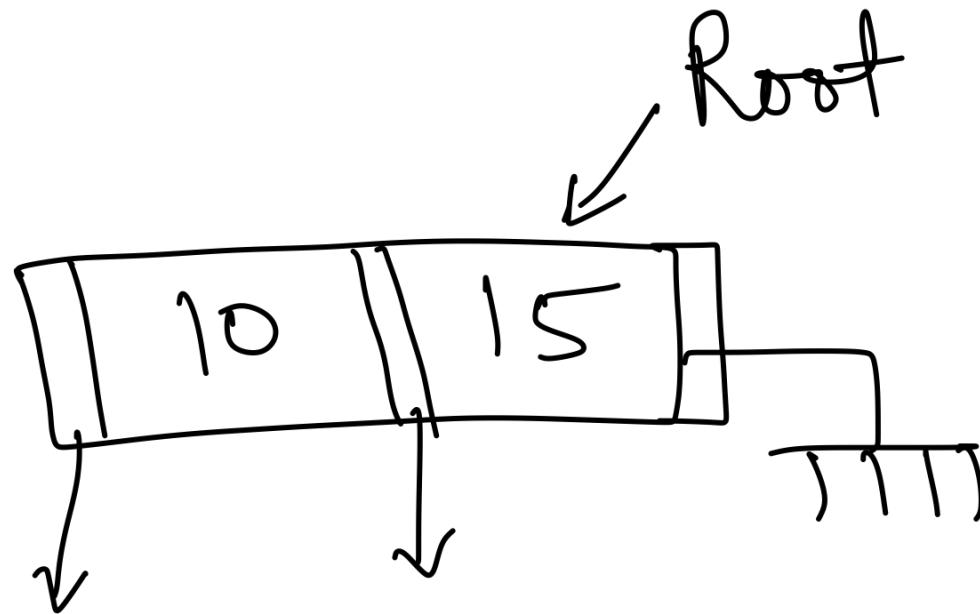


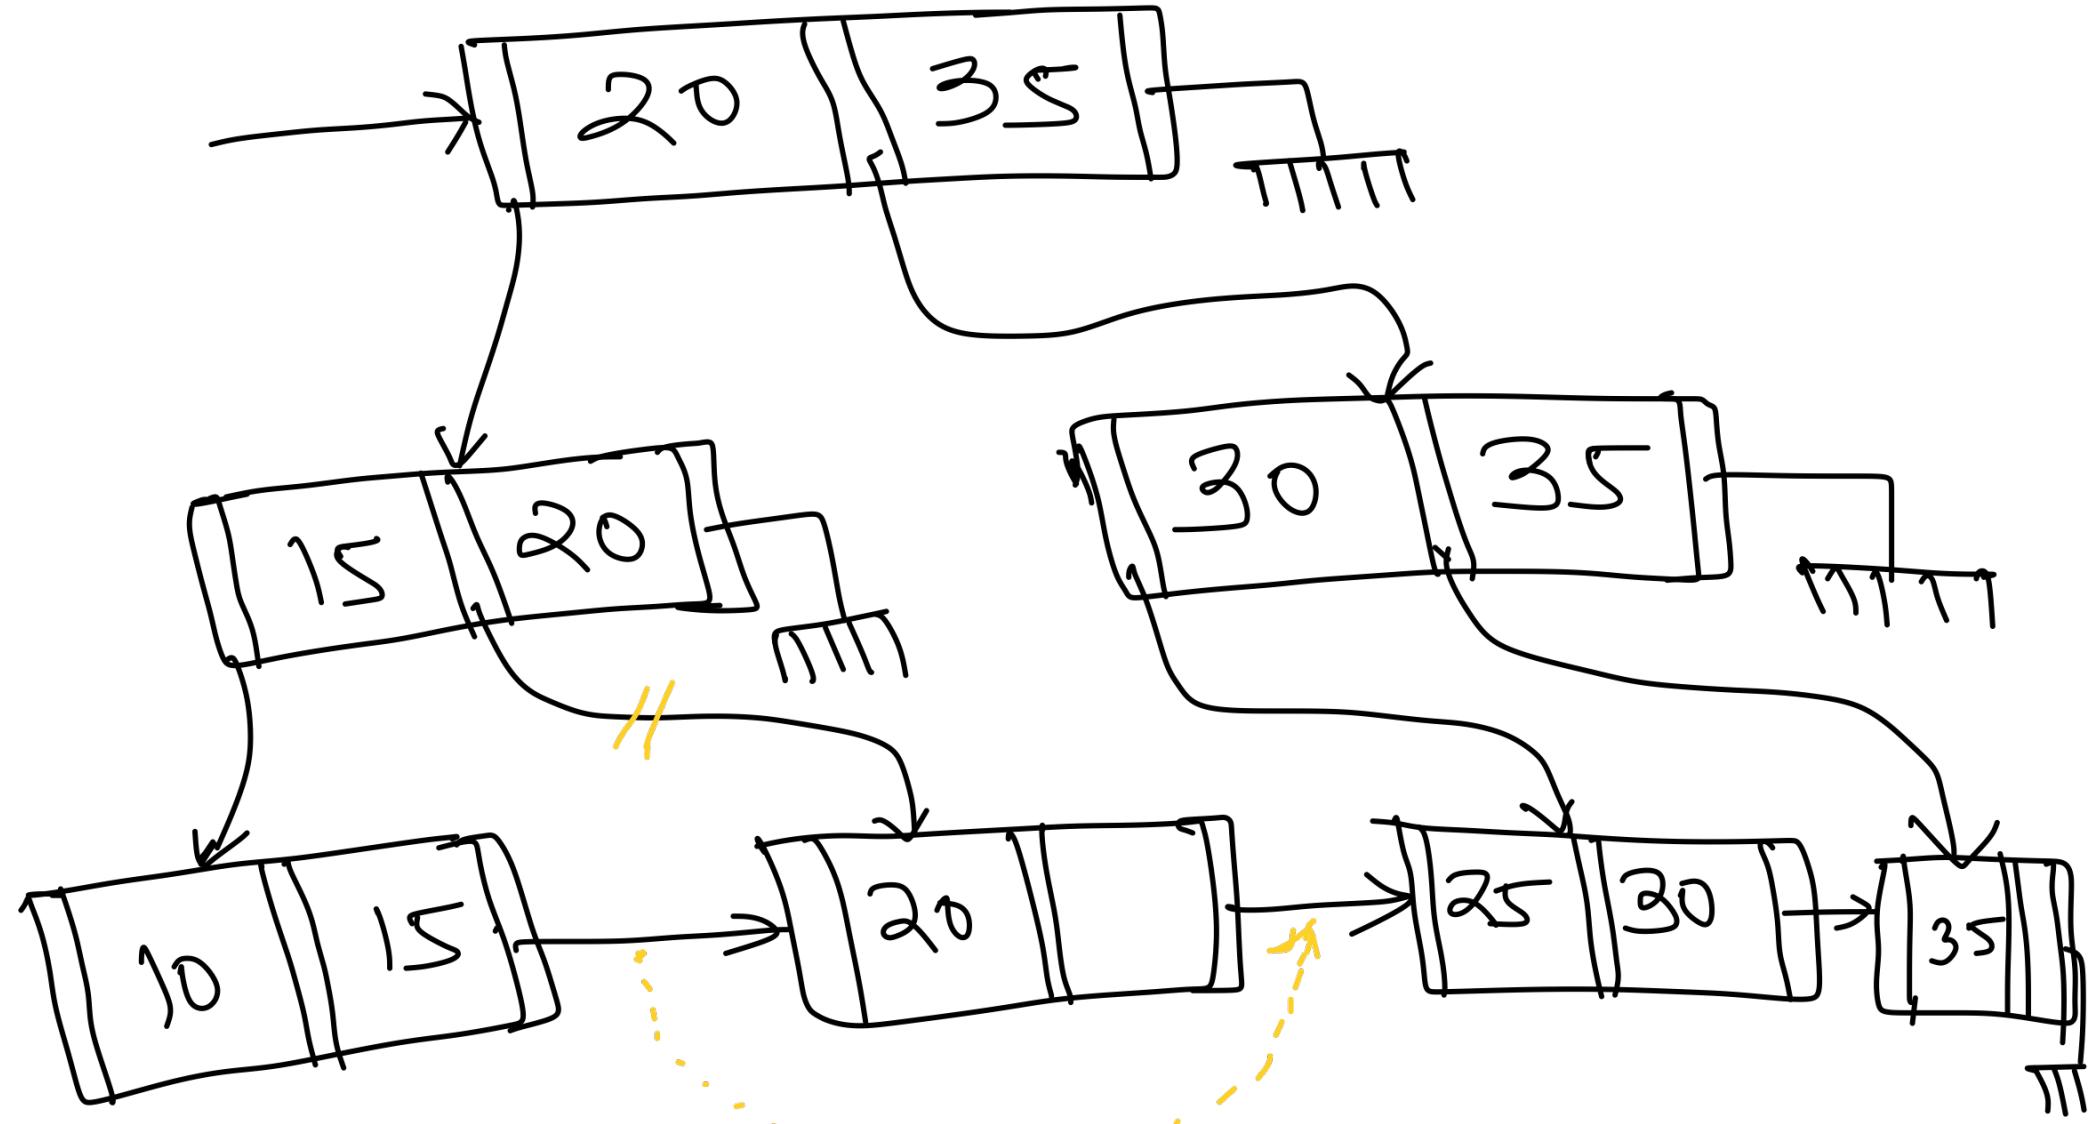
Delete 20





Delete 20

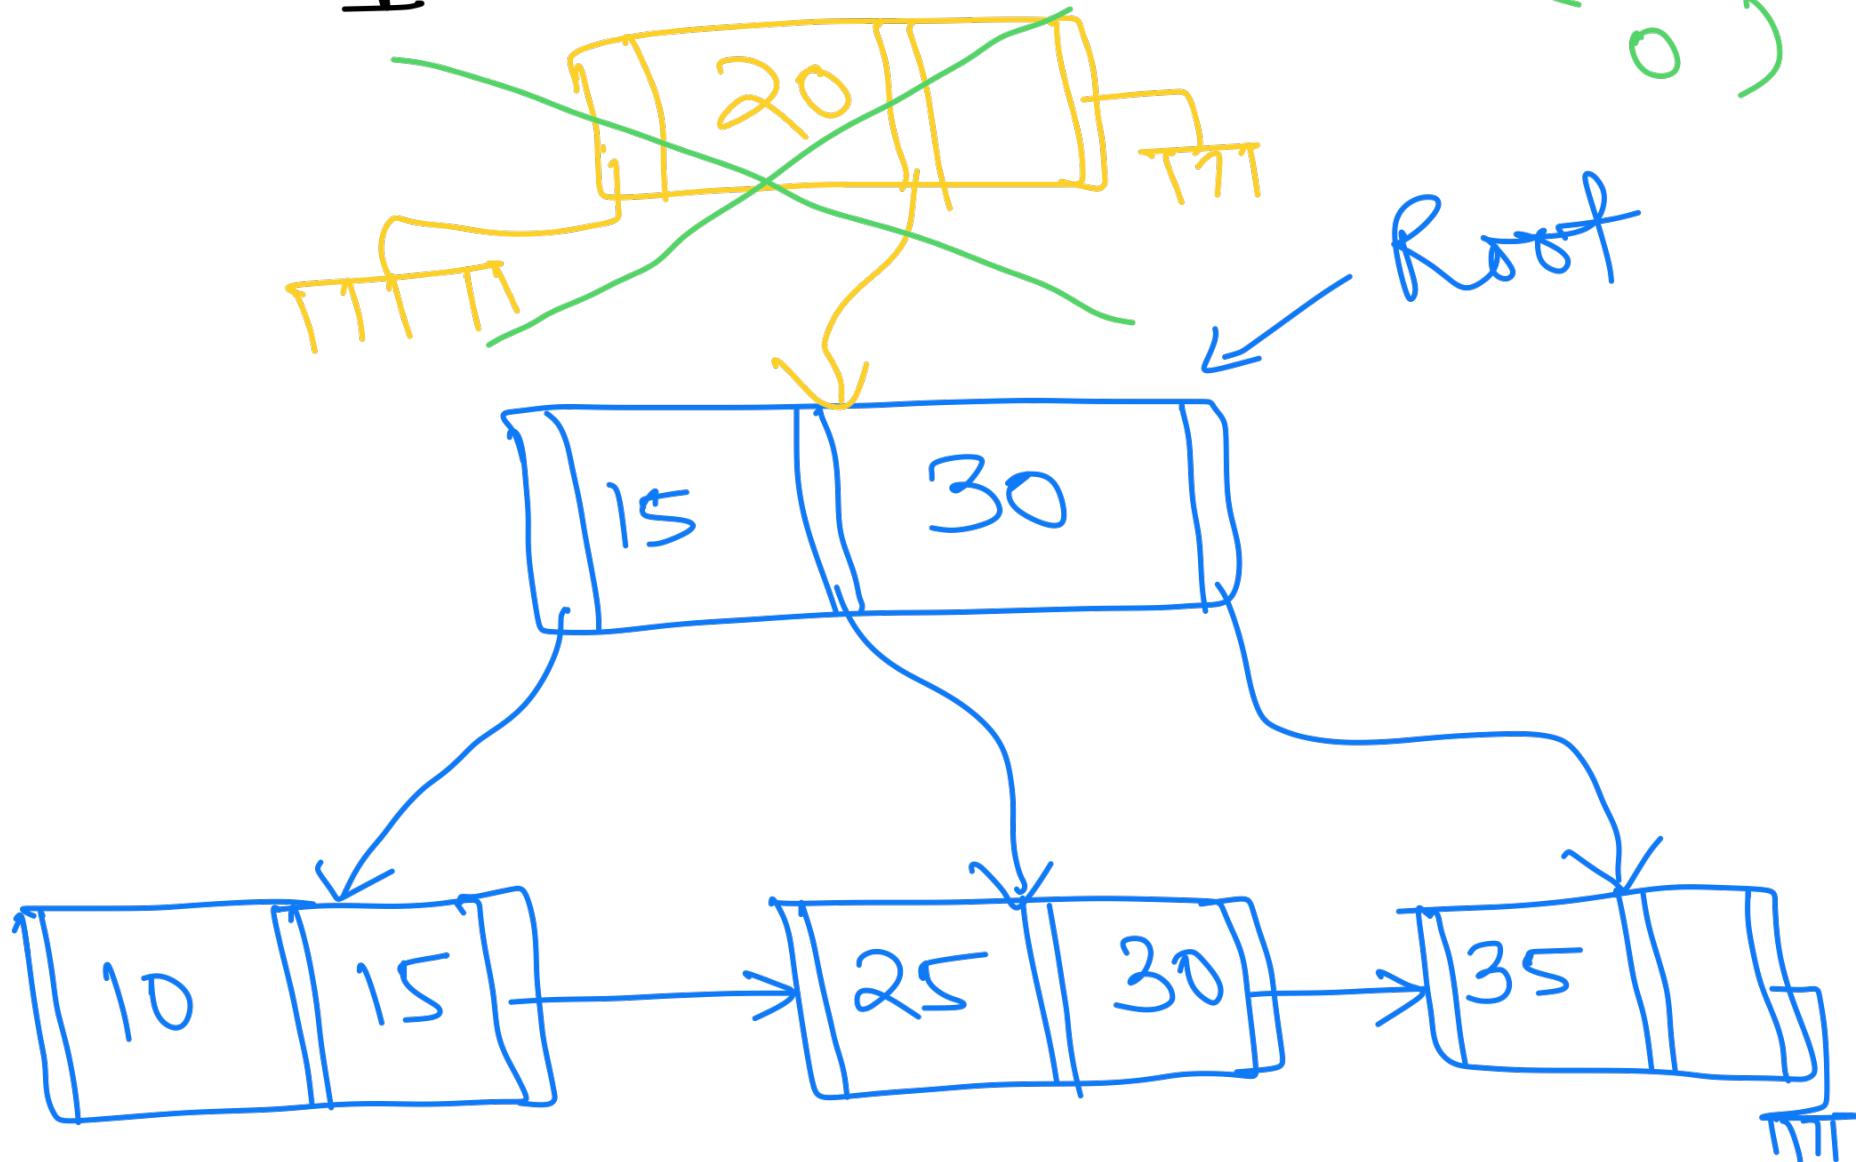


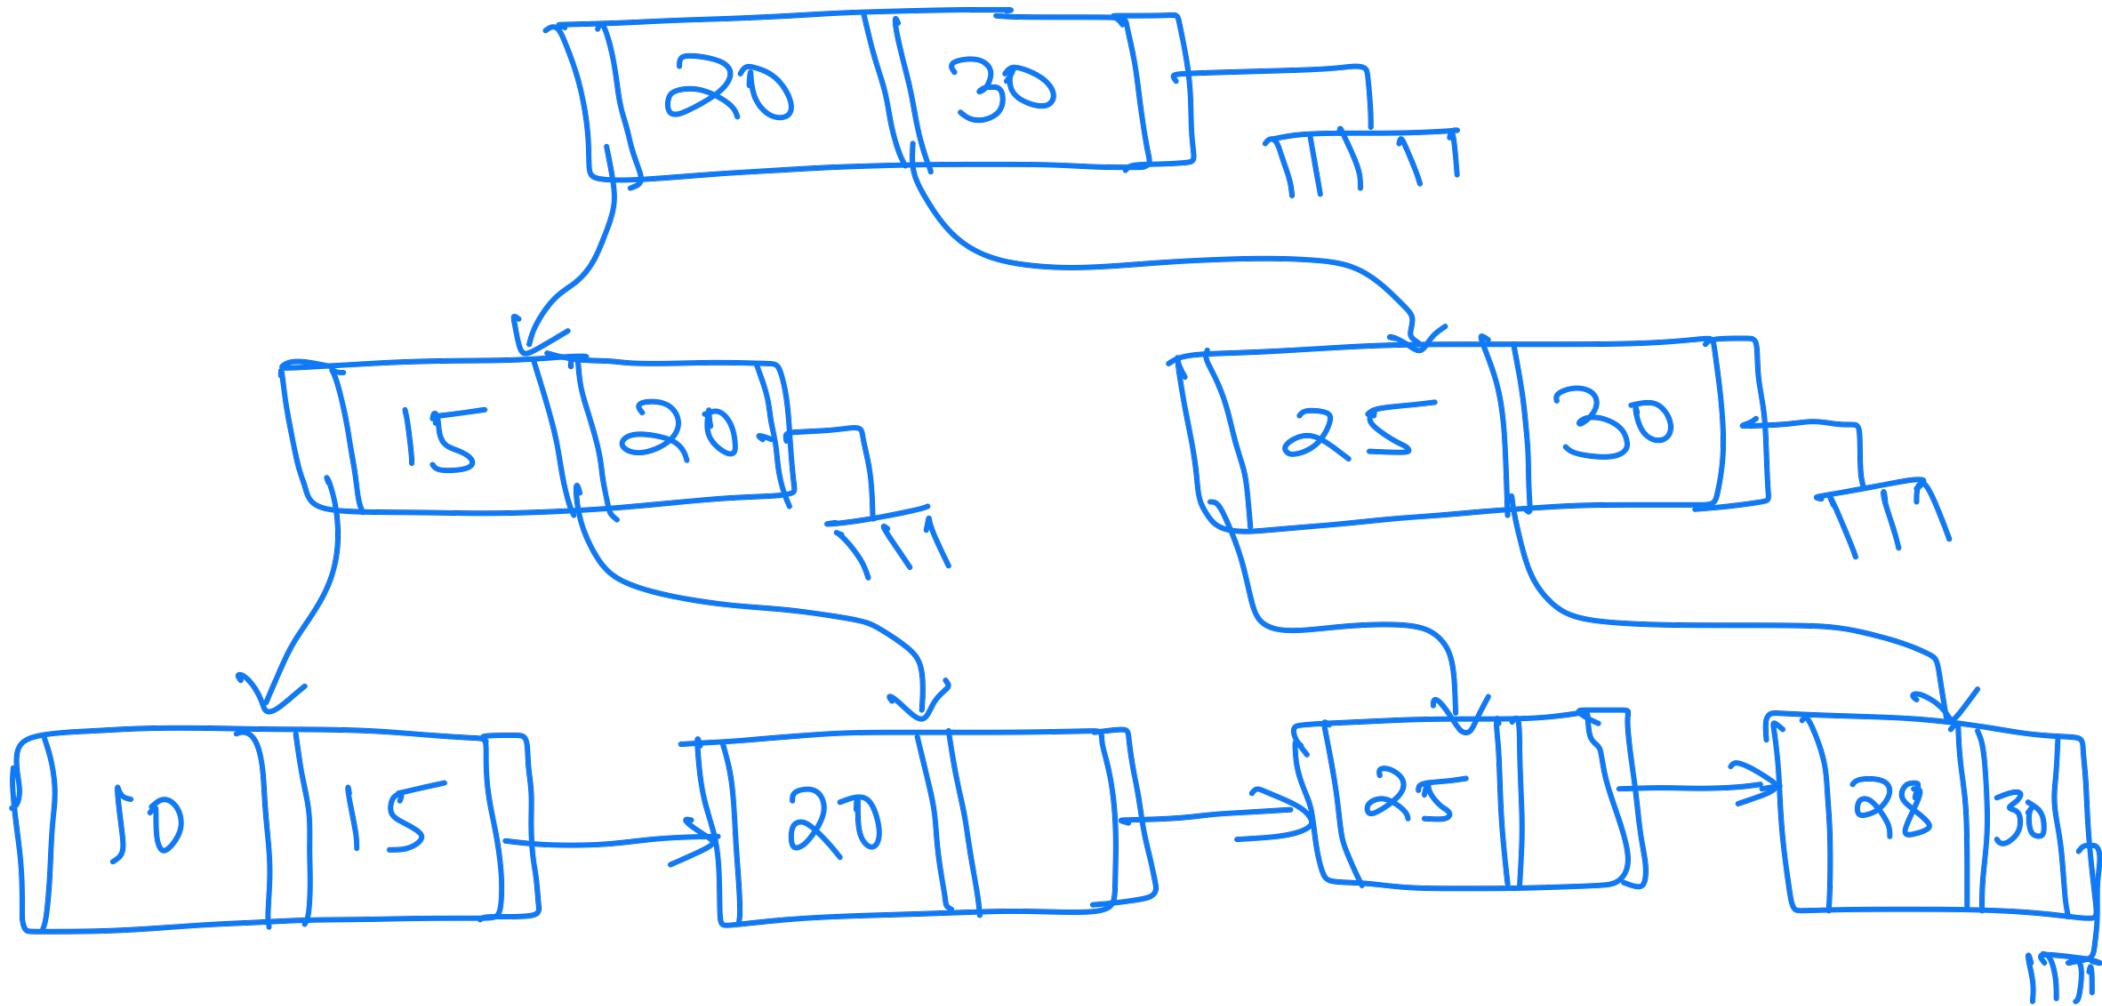


Delete 20

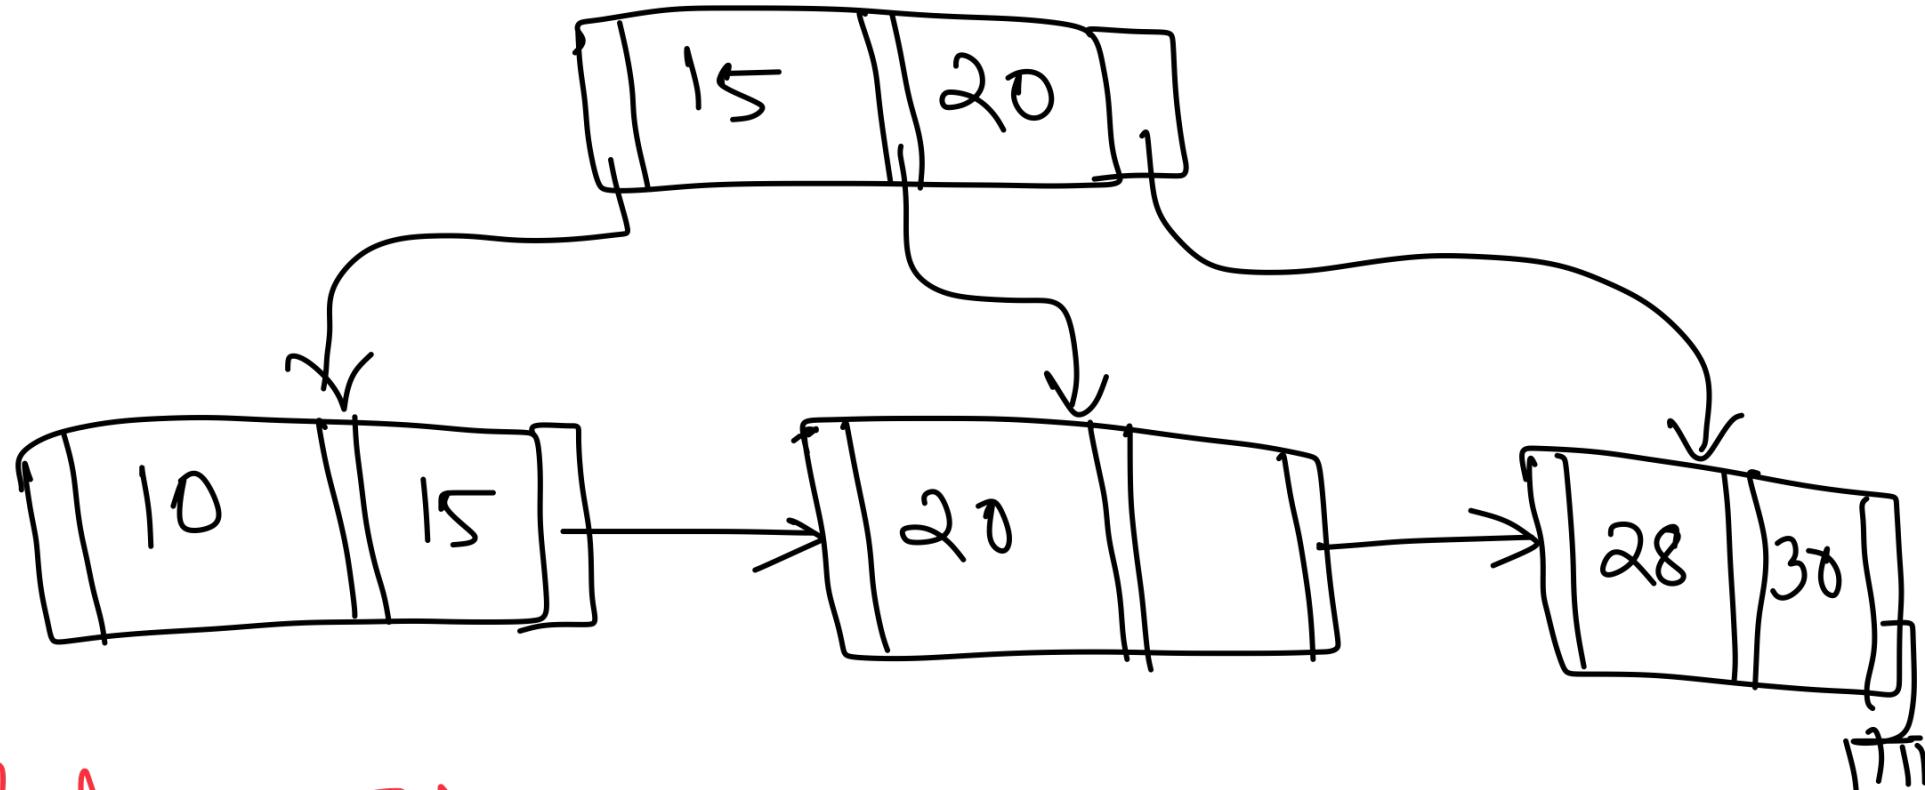
Case 1 : 1 3 → Redistribute 2 2

Case 2 : 1 2 → 11 03  
(Delete 0)





Delete 25



## Delete : Steps

- (i) Find the element
- (ii) Delete leaf <sup>node</sup> if required.
- (iii) Delete Internal node if 11.

## # Hashing: Hashed Index

$f(A) \rightarrow \text{Value}$

